

FOUNDATIONS OF COMPUTING 2

Introduction

- Comp. Sci. is the science of **automatic processing of information**.
- Comp. Sci. has its own **fundamental questions**:
 - What is an **algorithm**?
 - What are the **fundamental capabilities / limitations** of computers?
 - When should an algorithm be considered **practically feasible**?
 - What makes some problems **computationally hard**?

Turing Machines

- Basic features of an algorithm: (**informal**)
 - Clear, well-defined inputs and outputs
 - Organised structure
 - Reduces new problems to solved problems
 - eg. we can already solve division, squares, etc.
- These features are not always helpful though!
 - Example: Diophantine equations are well-defined, but do not readily reduce to solved problems / operations.
- Some problems (eg. Diophantine equations) are "solved" by proving that **there is no algorithmic solution**.
 - This "**non-existence**" statement cannot be made on the above, informal definition of an algorithm, so a **formal definition** is required.
- Formal definition of algorithms: **Turing machine**.
- A Turing machine has two parts:
 - A **tape**
 - A **finite control unit**

TM: Tape

- A sequence of cells, infinite in both directions.
- Each cell contains a symbol from a finite alphabet.
- There is a tape head that can:
 - Read the current symbol
 - Overwrite the current symbol
 - Move one cell left (L), one cell right (R), or stay at the current position (S)

TM: Finite State Control Unit

- Depending entirely on the current state and the current symbol in that state:
 1. Write a symbol into the current state
 2. Move according to L, R or S
 3. Enter and read the next state.

Decidability and Undecidability

- If we start a TM M on an input w , there are three possible outcomes:
 - M halts in an accepting state.
 - M halts in a rejecting state.
 - M loops and runs forever.
- A language is Turing-decidable if some TM decides it.
 - That is, some TM halts on all inputs.

The Acceptance Problem

- Testing whether a particular DFA accepts a string.

$$A_{DFA} = \{ \langle B, w \rangle \mid B \text{ is a DFA that accepts } w \}$$

- Theorem: A_{DFA} is a decidable language.
- Proof: Construct a TM that decides A_{DFA} :
 $M =$ "On input $\langle B, w \rangle$, simulate B on input w ; if simulation ends in an accepting state, accept, otherwise reject."

Testing Equality

- Do two DFAs recognize the same language?

$$EQ_{DFA} = \{ \langle A, B \rangle \mid A \text{ and } B \text{ are DFAs, } L(A) = L(B) \}$$

Theorem: EQ_{DFA} is a decidable language

- Proof: We use the previous acceptance proof.
 - Construct a new DFA C from A and B , where C only accepts strings accepted by $A \text{ XOR } B$.
 - If $L(A) = L(B)$ then $L(C) = \emptyset$.

$$L(C) = (L(A) \cap \overline{L(B)}) \cup (\overline{L(A)} \cap L(B))$$

The Halting Problem

- Some problems are algorithmically unsolvable
 - o They are undecidable.
- Test whether a given TM accepts a string w .

$$A_{TM} = \{ \langle M, w \rangle \mid M \text{ is a TM and } M \text{ accepts } w \}$$

Theorem: A_{TM} is undecidable *

Proof: The following TM U recognises A_{TM} :

U = "On input $\langle M, w \rangle$, simulate M on input w ; if M ever enters an accepting state, accept; if M enters a rejecting state, reject."

* In proving this theory, we show that requiring a TM to halt on all inputs restricts the kinds of languages it can recognise.
∴ A general TM is more powerful than one that halts on all inputs.

- Notes:
- U loops on input $\langle M, w \rangle$ if M loops on input w .
 - If there was a way to determine that M is non-halting on w , U could reject.
 - A_{TM} is called the halting problem.
 - U is called the Universal TM, because it can simulate any other TM from its description.

A Proof Idea:

- Assume A_{TM} is decidable and H is a TM that decides it.

$$H(\langle M, w \rangle) = \begin{cases} \text{accept} & \text{if } M \text{ accepts } w \\ \text{reject} & \text{if } M \text{ does not accept } w \end{cases}$$

- Construct a second TM D that uses H as a sub-routine.

D = "On input $\langle M \rangle$, run H on input $\langle M, \langle M \rangle \rangle$ and return the opposite of what H outputs"

$$D(\langle M \rangle) = \begin{cases} \text{accept} & \text{if } H \text{ does not accept } \langle M \rangle \\ \text{reject} & \text{if } H \text{ accepts } \langle M \rangle \end{cases}$$

- If we now run D on input $\langle D \rangle$, we have:

- accept if D does not accept $\langle D \rangle$
- reject if D accepts $\langle D \rangle$

- This is an **obvious contradiction**.

TM Variants

- Multiple tapes
- Multiple heads
- Nondeterministic TMs
- Random Access TMs

- Two-dimensional tapes
- etc...

- **None of these add power**: we show equivalence by simulating one on the other.

Algorithms

- A well-defined, finite set of instructions to perform a calculation or solve a problem.
- Of an algo, we can ask... is it **correct**?
is it **efficient**? → can use Big-Oh
is it **optimal**?

Example: Fibonacci Numbers

- Simple, exponentially recursive solution:

```
fib(n) {  
    if (n == 0) return 0;  
    if (n == 1) return 1;  
    return fib(n-1) + fib(n-2);  
}
```

- Correct
- Very inefficient
- Runs in exponential time

- Dynamic Programming solution:

```
fib(n) {  
    create array f[0...n]  
    f[0] = 0;  
    f[1] = 1;  
    for i = 2 ... n:  
        f[i] = f[i-1] + f[i-2];  
    return f[n];  
}
```

- Correct
- More efficient
- Runs in linear time

- Iterative solution:

```
fib(n) {  
    if (n <= 1) return n;  
    x = 0; y = 0; z = 1;  
  
    for (i = 1; i < n; ++i) {  
        x = y;  
        y = z;  
        z = x + y;  
    }  
    return z;  
}
```

- Correct
- Uses less memory than prev. solution
- Runs in linear time

- Matrix solution:

$$\begin{pmatrix} F_{n-1} \\ F_n \end{pmatrix} = \begin{pmatrix} 0 & 1 \\ 1 & 1 \end{pmatrix}^{(n-1)} \cdot \begin{pmatrix} 0 \\ 1 \end{pmatrix}$$

- Correct
- Efficiency depends on exponentiation method

Example: Exponentiation

- Basic method:

$$a^n = \underbrace{a \times a \times \dots \times a}_n \quad - \text{Runs in linear time}$$

- By squaring:

$$a^n = \begin{cases} 1 & \text{if } n=0 \\ a & \text{if } n=1 \\ (a^2)^{\frac{n-1}{2}} & \text{if } n \text{ is even} \\ (a^2)^{\frac{n-1}{2}} \cdot a & \text{if } n \text{ is odd} \end{cases} \quad - \text{Runs in } \log_2 n \text{ time} \\ - \text{Fast!}$$

Primes

- Used in various areas of **cryptography**.
- Huge numbers must be generated - upwards of 500 digits.

The Sieve of Eratosthenes (250 BC)

- Input: integer n
- Output: a list of all primes $\leq n$

- Process:

```
P = Ø;  
L = [2, 3, 4, ..., n];  
while L ≠ [] do:  
    a = head(L);  
    P = P ∪ a;  
    delete all multiples of a from L;  
return P;
```

Prime Test: Naive Division Method

- Input: integer n
- Output: true if n is prime; false otherwise

- Process:

```
if (n == 2) return true;  
if (n % 2 == 0) return false;  
for (int i = 1; i <= ⌊√n/2⌋; ++i):  
    if (n % (2i+1) == 0) return false;  
return true;
```

- Perform $O(\sqrt{n}) = O(n^{1/2}) = O(2^{(\log_2 n)^{1/2}})$ division

• i.e. Exponential in $\log_2 n$

- Very slow - useless in practice

- Randomised algorithms can be a lot faster.

Fermat Little Theorem

- Let p be a prime number.

- For all $a \neq 0$, $a^{p-1} \pmod p \equiv 1$

- Eg. $p = 7$
 $a = 4$

$$a^6 = 4096$$

$$4096 \pmod 7 = 1.$$

Also stated as

$$a^{p-1} \equiv 1 \pmod p$$

Quotient-Remainder Theorem

- Let a be an integer, and b be a positive integer.

- There exist unique integers q and r such that:

- $0 \leq r < b$

- $a = qb + r$

- q is the quotient: $a \div b$

$$7 \div 2 = 3$$

$$7 \pmod 2 = 1$$

- r is the remainder: $a \pmod b$

- For $a > 0$ and $b > 0$, we define:

$$a \pmod b = \begin{cases} a & \text{if } a < b \\ (a-b) \pmod b & \text{if } a \geq b \end{cases}$$

Modular Arithmetic

- Two or more integers may be congruent, modulo a number N .

o eg. $7 \equiv 19 \pmod 2$

because... $7 \pmod 2 = 1$

$$19 \pmod 2 = 1.$$

- If $r = a \pmod N$, then $r \in \{0, 1, \dots, N-1\}$

- We say two numbers congruent modulo n : $a \equiv b \pmod N$

- Let $a_1 \equiv b_1 \pmod{n}$ and $a_2 \equiv b_2 \pmod{n}$, then ...
 - $a_1 + a_2 \equiv b_1 + b_2 \pmod{n}$
 - $a_1 - a_2 \equiv b_1 - b_2 \pmod{n}$
 - $a_1 a_2 \equiv b_1 b_2 \pmod{n}$

Randomised Primality Testing

- From Fermat Little Theorem we know that if p is prime then:
$$\forall a \in \{1, \dots, p-1\} : a^{p-1} \equiv 1 \pmod{p}$$
- Process:
 - Given a number n
 - Randomly select $a \in \{1, \dots, n-1\}$
 - Check whether $a^{n-1} \equiv 1 \pmod{n}$ holds
 - If not, then we know n is not prime
 - We say a is a witness
 - $a^{n-1} \equiv 1 \pmod{n}$ does not guarantee n is prime
- Checking $a^{n-1} \equiv 1 \pmod{n}$ can be done in $O(\log n)$.
- So the above algo can run in $O(n \log n)$.

Carmichael Numbers

- The Fermat-Test can only prove with certainty that a number is not prime.
 - A positive result is not guaranteed to be correct.
 - CNs are numbers that are not prime, but have very few witnesses.
 - They have this property:
- $$\forall (a \text{ with } \gcd(a, n) = 1) : a^{n-1} \equiv 1 \pmod{n}$$

\gcd : greatest common divisor

Non-Trivial Square Roots

- Let p be prime and $a \in \{1, \dots, p-1\}$
- For $a^2 \equiv 1 \pmod{p}$ we only have $a = 1$ and $a = p-1$
- If there is an $a \in \{2, \dots, n-2\}$ such that $a^2 \equiv 1 \pmod{n}$, then a is a non-trivial square root modulo n

- Eg. Let $a = 6$, $n = 35$

$$6 \in \{2, \dots, 33\}. \quad 6^2 \equiv 1 \pmod{35}$$

$\therefore 6$ is a non-trivial square root modulo 35

$\therefore 35$ is not prime

- This can be used to improve the randomised algorithm.

Modular Exponentiation

- From the last lecture:

$$a^n = \begin{cases} a^{\frac{n}{2}} \times a^{\frac{n}{2}} & \text{if } n \text{ is even} \\ a^{\frac{n-1}{2}} \times a^{\frac{n-1}{2}} \times a & \text{if } n \text{ is odd} \end{cases}$$

- This algo will compute an exponentiation mod n , and give an indication of primality.

o Input: integers a, p, n

o Output: $r = a^p \pmod{n}$, prime = false if there is an NTSR, true otherwise

o Process: prime = true;

power(a, p, n): if ($p == 0$) { $r = 1$;}

else {

$x = \text{power}(a, \text{int}(p/2), n);$

$r = (x \cdot x) \pmod{n};$

if ($r == 1$ and $x \neq 1$ and $x \neq n-1$) prime = false;

}

if (p is odd) $r = r \cdot a;$

return $r, \text{prime};$

Miller-Rabin Algorithm

- This combines the Fermat-Test with the test for NTSRs.

- Input: integer n

- Output: $\text{false} \rightarrow n \text{ is not prime}$

$\text{true} \rightarrow n \text{ is probably prime}$

- Process:

```
a = rand(2, n-1);  
r, prime = power(a, n-1, n);  
if (r ≠ 1 or !prime) return false;  
return true;
```

- If n is not prime, there are at least $(n-1)/2$ numbers $a \in \{2, \dots, n-1\}$ that are witnesses for n being a non-prime in the MR algo test.

• The probability of the MR test giving true for a non-prime is $\frac{1}{2}$.

• Run the test 50 times, and this probability is $\frac{1}{2^{50}}$ or 2^{-50} .

Public key Encryption

Private keys:

- Alice and Bob have a shared private key, k .

- Alice computes $V(\text{Msg}, k) = C$

- C is sent to Bob

- Bob computes $V^{-1}(C, k) = \text{Msg}$

- Problems: • how do they share k ?

• multiple people $\rightarrow n(n-1)/2$ keys!

Public keys:

- Each party has two keys: public and private (secret).

P_A = Alice's public key

S_A = Alice's private key

- Bob gets P_A and encrypts his message: $P_A(\text{Msg}) = C$

- C is sent to Alice

- Alice uses her private key: $S_A(C) = \text{Msg}$

Key Properties

- Let D be the set of all possible messages (all finite binary strings).
- Each party has two bijective functions:

$$P_A: D \rightarrow D$$

$$S_A: D \rightarrow D$$

that satisfy...

- P_A and S_A can be efficiently computed.

$$\bullet \forall m \in D: P_A(S_A(m)) = m$$

$$\bullet \forall m \in D: S_A(P_A(m)) = m$$

- S_A is hard to get from P_A .

- The third property is used for sending encrypted messages.
- The second property can be used to sign a file.

Digital Signatures

- Alice computes $\sigma = S_A(f)$ and sends f to Bob
- Anyone can prove with $P_A(\sigma) = P_A(S_A(f)) = f$ that the file really was sent by Alice
- Note: σ depends on f and is only valid for one document.
- Note: it is usually sufficient to compute $\sigma = S_h(f)$ where h is some hash function.

RSA Encryption Algorithm

gcd = greatest common divisor

- To generate keys:

- o Pick two large primes p and q

- o Let

$$n = p \times q$$

$$\phi = (p-1) \times (q-1)$$

e = a small number w/ $\text{gcd}(e, \phi) = 1$.

- o Pick d such that $d \times e \pmod{\phi} = 1$ (mult. inverse)

- o Publish $P(e, n)$ as the public key

- o keep $S(d, n)$ as the private key.

- Sender encrypts m : $P(m) = M^e \pmod{n} = c$

- Receiver decrypts m : $S(c) = C^d \pmod{n} = m$

Computing GCD

- $\text{gcd}(a, b)$ = the largest integer divisor of a and b .

- Computing w/ factorisation:

- o Factorise a and b

- o $\text{gcd}(a, b)$ = product of common factors.

- But factorisation is computationally hard.

- Euclid's Algorithm:

$$\text{gcd}(a, 0) = a$$

$$\text{gcd}(a, b) = \text{gcd}(b, a \pmod{b}) \quad b > 0$$

- Note: if $\text{gcd}(a, b) = d$, there exist integers m and n such that:

$$ma + nb = d$$

-Proof of Euclid's Algorithm

Let $c = \gcd(a, b)$ and $d = \gcd(b, r)$.

Note that $a = qb + r$

from $a = qb + r$, and $d | b, d | r \rightarrow d | a$

because $d | a, d | b$ and $c = \gcd(a, b) \rightarrow d | c$

$\therefore c > d$

From $r = a - qb$ and $c | a, c | b \rightarrow c | r$

because $c | b, c | r$ and $d = \gcd(b, r) \rightarrow c | d$

$\therefore d > c$

Therefore, $c = d$

Computing Multiplicative Inverse.

Say we have $e=5$ and $s=6$.

We need to pick a number d such that $5d \pmod{6} = 1$.

Suitable options: 5, 11, 17, etc...

$$0 < d \quad (d \text{ less than } 6) \pmod{6} = (d, 0) \pmod{6}$$

$$d = dn + DM$$

Full example:

- Pick $p = 13$ and $q = 7$
- $n = 91$
- $s = 72$
- $e = 5$
- $d = 29$
- $P(s, 91)$
- $S(29, 91)$

$$\begin{aligned}e \text{ Proof: } & \gcd(s, 72) \\&= \gcd(5, 2) \\&= \gcd(2, 1) \\&= \gcd(1, 1) \\&= \gcd(1, 0) \\&= 1\end{aligned}$$

$$d \text{ Proof: } 5 \times 29 = 145$$

$$145 \bmod 72 = 1$$

- Encrypt:

$$\text{Message} = 6$$

$$\text{Ciphertext} = 6^5 \bmod 91 = 41$$

- Decrypt:

$$\text{Message} = 41^{29} \bmod 91 = 6$$

Correctness of RSA

- Does RSA satisfy the three conditions from earlier?

1. Can $P(m)$ and $S(m)$ be efficiently computed?

- Yes! M^e and M^d can be computed via exponentiation modulo n in $O(\log e) = O(\log d) = O(\log n)$.

2. Are $P(m)$ and $S(m)$ mutually inverse?

- p and q are prime, so Fermat's Little Theorem tells us:

$$M^{p-1} = 1 \bmod p$$

$$M^{q-1} = 1 \bmod q$$

- The Chinese Remainder Theorem tells us:

$$M^{(p-1)(q-1)} = 1 \pmod{p} \quad \text{and} \quad M^{(p-1)(q-1)} = 1 \pmod{q}$$

- So... $S(P(m)) = (m^e)^d \pmod{n}$

$$= M^d \pmod{n}$$

$$= M^{1 + e(p-1)(q-1)} \pmod{n}$$

$$= m \pmod{n}$$

- Because we picked $e \times d = 1 \pmod{s}$, and $s = (p-1)(q-1)$.

- $\therefore e \times d = 1 + r \times s$ for some r ,

- The same can prove the inverse, $P(S(m))$.

- So yes, property two is satisfied.

3. Is $S()$ hard to guess from $P()$?

- Computing d is easy if we know p and q .

- But we only know e and n , which $= pq$, but we don't know them.

- To "crack" a message, we must factorise n .

- This is very hard! So yes, this passes.

Graphs

- A graph G is defined by the triple (V, E)
 - V = a non-empty set of vertices
 - E = a set of edges (pairs of vertices)
- There are several types of graphs:
 - Undirected graphs - edges are unordered pairs (v, w)
 - Directed graphs - edges are ordered pairs $\langle v, w \rangle$
 - Multi-graphs
 - multiple edges between the same two vertices, and/or "looping" edges.
 - Multi-graphs can be directed or undirected.

Terms

- Adjacent vertices: v is adj. to w iff $\exists (v, w) \in E$
- Degree: $\deg(v) = \# \text{ of vertices adjacent to } v$.
- Path: a sequence of vertices v_0, v_1, \dots, v_k with $\exists (v_i, v_{i+1}) \in E$ for all $0 \leq i < k$
- Simple Path: path with no repeated vertices
- Length of a path = # of edges
- Cycle: a "closed" path with $v_0 = v_k$
- Simple cycle: a "closed" simple path.

Connectivity

- Connected graph: for any two vertices, a path can be found that joins them
- Subgraph: $H = (V', E')$ is a subgraph of $G = (V, E)$ if $V' \subseteq V$ and $\forall e \in E', e = (v, w) : v, w \in V'$ and $(v, w) \in E$
- Spanning subgraph: a subgraph where $V' = V$
- Induced subgraph: a subgraph where no edges have been "lost" from the vertices that remain.
- Connected component: a maximal induced subgraph that is connected.

Complete Graphs

"Most
Connected"

- A complete graph has an edge connecting each pair of vertices.
- Complete graph of n vertices = k_n
- aka. "Cliques"

Tree Graphs

"Least
Connected"

- No cycles
- One path between any pair of vertices
- Removing any edge disconnects the graph

Edges in Complete Graphs

$$- k_1 = 0, \quad k_2 = 1, \quad k_3 = 3, \quad \dots \quad k_n = \frac{n(n-1)}{2}$$

- Proof: $k_{n+1} = k_n + n$ (from k_n we have to add n edges to draw k_{n+1})

$$k_{n+1} = \frac{n(n-1)}{2} + n = \frac{n(n-1) + 2n}{2} = \frac{n(n+1)}{2}$$

Edges in Tree Graphs

- A tree graph with n vertices has $n-1$ edges.

Spanning Trees

- A spanning tree of a graph G is a spanning subgraph that is also a tree.

- A complete graph k_n has n^{n-2} distinct spanning trees.

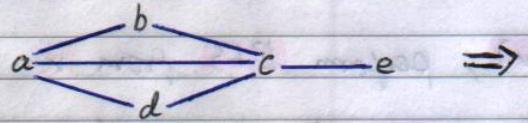
- Spanning trees can be found by starting at a node and applying depth-first or breadth-first searches, recording the edges that lead to unvisited nodes.

Graph Representations

- Adjacency List

- For all $v \in V$, $L(v)$ lists all w such that $(v, w) \in E$.

Eg.

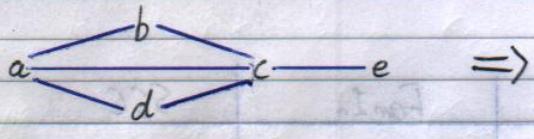


v	$L(v)$
a	b c d
b	a c
c	a b d e
d	a c
e	c

- Adjacency Matrix

- An $n \times n$ matrix such that $A[v, w] = 1$ means $(v, w) \in E$.

Eg.



	a	b	c	d	e
a	0	1	1	1	0
b	1	0	1	0	0
c	1	1	0	1	1
d	1	0	1	0	0
e	0	0	1	0	0

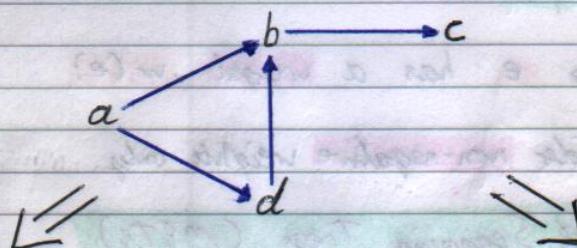
Symmetric
for a simple
graph.

Digraphs

- Edges are ordered pairs: $(v, w) \Leftrightarrow v$ is adj. to w

- (v, w) does not imply (w, v)

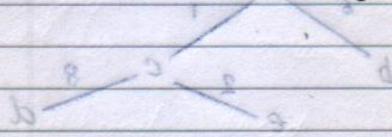
Eg.



v	$L(v)$
a	b d
b	c
c	
d	b

	a	b	c	d
a	0	1	0	1
b	0	0	1	0
c	0	0	0	0
d	0	1	0	0

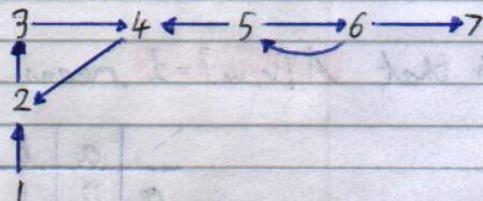
$$8 + 2 + 1 + 3 = 14$$



- $\text{FanOut}(v)$: set of all vertices reachable from v (incl. v)
- $\text{FanIn}(v)$: set of all vertices from which v is reachable (incl. v)
- $\text{SCC}(v)$: strongly connected component containing v

$$\text{sec}(v) = \text{FanIn}(v) \cap \text{FanOut}(v)$$

- To find $\text{FanOut}(v)$, perform BFS from v
- To find $\text{FanIn}(v)$, reverse all edges and perform BFS from v
- Eg:



Node	FanOut	FanIn	SSC
1	1 2 3 4	1	1
2	2 3 4	2 3 4 5 6 1	2 3 4
3	3 4 2	3 2 4 5 6 1	2 3 4
4	4 2 3	4 1 2 3 5 6	2 3 4
5	5 6 7 4 2 3	5 6	5 6
6	6 5 7 4 2 3	6 5	5 6
7	7	7 6 5	7

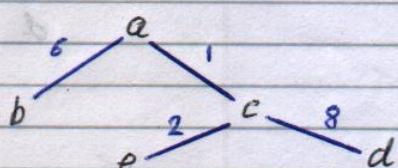
Weighted Graphs

- Each edge e has a weight $w(e)$
- We consider non-negative weights only.

Minimum Cost Spanning Trees (MSTs)

- An MST for the graph G is a spanning tree of G such that the combined sum of all edges is minimized.

- Eg. $T_a =$



$$w(T_a) = 6 + 1 + 2 + 8 \\ = 17$$

- Prim's Algorithm for finding MSTs

◦ Input: graph G with vertices V and edges E

◦ Output: T , a subset of E inducing MST

◦ Process:

$$S = \{ \}$$

(set of seen vertices)

Start at some vertex v

$$S = \{ v \}$$

Until T spans G , repeat:

Add the cheapest edge e joining an unvisited vertex w to T

$$T = T \cup e$$

$$S = S \cup w$$

Return T

- Kruskal's Algorithm for finding MSTs (Growing Forests)

◦ Grows a forest of trees and finds the cheapest edge to connect two disjoint trees.

◦ Input: graph G with vertices V and edges E

◦ Output: T , a subset of E inducing MST

◦ Process:

$$T = \{ \}$$

(set of edges belonging to MST)

Start w/ all vertices and no edges

Sort all edges in order by weight

$$T = \{ e_i \}$$

Until T spans G , repeat:

Does the next cheapest edge (v, w) join two distinct connected components of T ?

$$T = T \cup \{ (v, w) \}$$

Return T

Single Source Shortest Path

- The shortest path P from nodes v to w is a path from v to w such that the cost of edges in that path is minimized.

- One approach to find this path is Dijkstra's Algorithm!

- Dijkstra's Algorithm

• We set $w(e) = \infty$ for all $e \notin E$

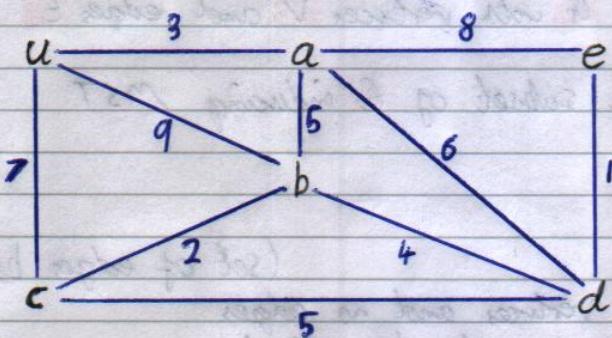
◦ Input: Coraph $G = (V, E)$ and vertex $u \in V$

◦ Output: shortest path from u to all other vertices

◦ Process:

Initial Phase { $S = \{u\}$ (set of seen vertices)
 $\ell(u) = 0$ (cost of reaching this vertex)
 for all $y \in V$ do $\ell(y) = w(u, y)$ *
 while $S \neq V$:
Looping Phase { Choose $x \in V - S$ with smallest $\ell(x)$ and add x to S
 for all $y \in V - S$:
 $\ell(y) = \min(\ell(y), \ell(x) + w(x, y))$
 return paths according to ℓ

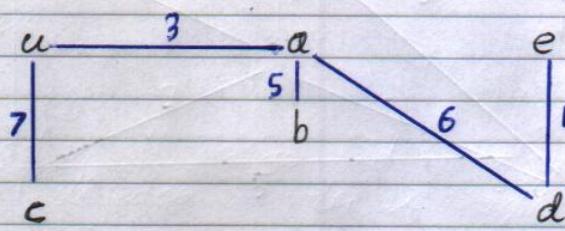
- Dijkstra Example Coraph:



Problem: Find the shortest path from u to all other nodes.

x	$V - S$	$l(y)$	$l(x) + w(x,y)$	min	from	edge
Initial Phase	u	a	3	-	(3)*	u u, a
		b	9	-	9	
		c	7	-	7	
		d	∞	-	∞	
		e	∞	-	∞	
Looping Phase	a	b	9	$3 + 5 = 8$	8	$u^{\Delta 1}$ u, c
		c	7	$3 + \infty = \infty$	(7)*	
		d	∞	$3 + 6 = 9$	9	
		e	∞	$3 + 8 = 11$	11	
	c	b	8	$7 + 2 = 9$	(8)*	$a^{\Delta 2}$ a, b
		d	9	$7 + 5 = 12$	9	
		e	11	$7 + \infty = \infty$	11	
	b	d	9	$8 + 4 = 12$	(9)*	$a^{\Delta 3}$ a, d
		e	11	$8 + \infty = \infty$	11	
	d	e	11	$9 + 1 = 10$	10*	$d^{\Delta 4}$ d, e

Result:



* = "choose an x with the lowest $l(x)$ "

$\Delta 1$: c was given 7 during the u round

$\Delta 2$: b was given 8 during the a round

$\Delta 3$: d was given 9 during the a round

$\Delta 4$: e was given 10 during the d round.