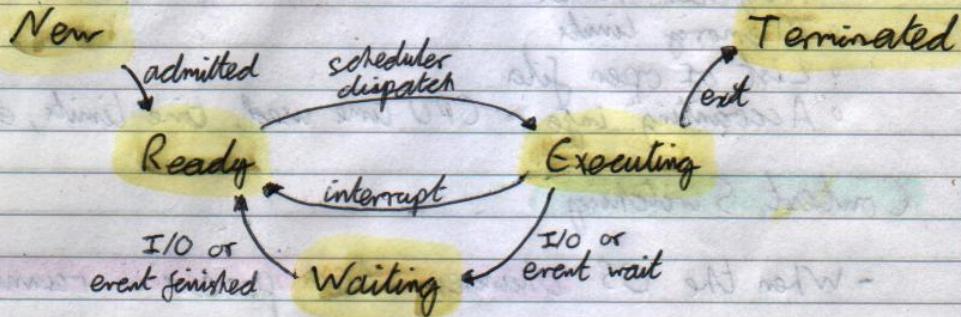


OPERATING SYSTEMS + CONCURRENCY

Processors

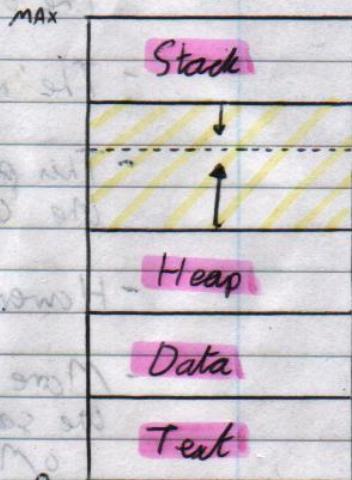
- A process is a program at some stage of execution.
- It has a program counter and variable values.
- A program is an unchanging set of instructions.

Process States



Processes in Memory

- **Stack:** temporary data like function parameters, return addresses and local variables.



- **Heap:** dynamically allocated memory (new objects); the memory the process uses.

- **Data:** global variables (e.g. static variables)

- **Text:** compiled binary for the program

- NB: register values / program counters are stored in registers when a process is executing and in the PCB when it is not.

Stack Overflow

→ Next page →

- The stack has a fixed size - overflow is often caused by a non-returning recursion.

Process Control Block (PCB)

- The PCB is maintained by the OS to represent state and scheduling information for processes.
- It is stored separately from processes.
- It contains...
 - Process state - ready, waiting, etc ...
 - Process number - a UID
 - Program counter
 - Register values
 - Memory limits
 - List of open files
 - Accounting info - CPU time used, time limits, etc...

Context Switching

- When the OS changes which process is running
- The "old" process has register values moved TO the PCB.
- The "new" process has register values restored FROM the PCB.
- This process is an overhead: it takes time, during which the CPU does no useful work.
- However, it allows multiple processes to "share" the CPU.
- More work gets done: CPU-bound processes can run at the same time as I/O-bound ones.
 - Most processes only require short bursts of CPU.

Process Scheduling

- An OS has several queues:
 - Job queue - all of the processes in the system
 - Ready queue - processes waiting for CPU time
 - Device queue - one per device; process waiting to use it
- Processes will move between queues during execution.

- There are two Phases of Scheduling:

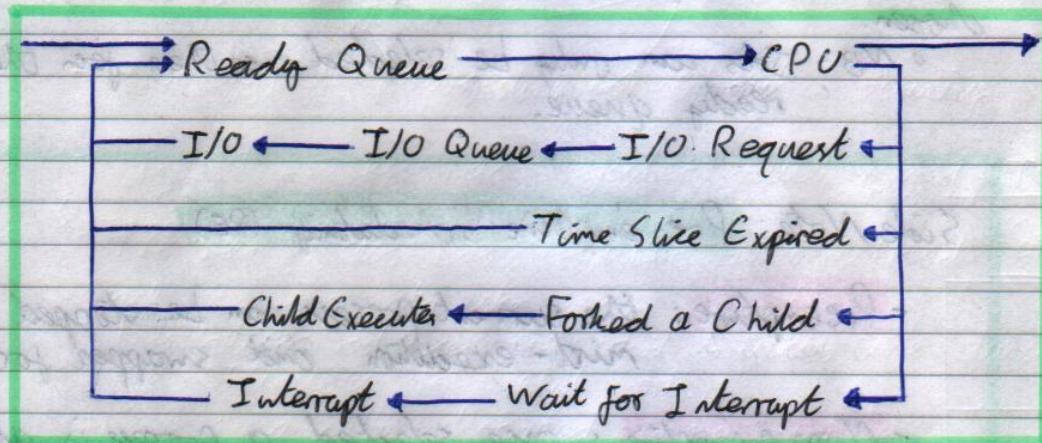
- o Long-term scheduler:

- Determines which processes join the ready queue.
 - Runs infrequently
 - Determines the degree of multiprogramming (how many processes can run at once).

- o Short-term scheduler

- Puts from the ready queue to run processes on the CPU
 - Runs frequently, so must be fast.

- Everything starts on the ready queue once added by the long-term scheduler, executes on the CPU for some time, then is moved to another queue:



- Good efficiency can be achieved by having multiple processes (I/O and CPU) at different stages in this system at once.
- Most processes do CPU work in short bursts, so we talk about "duration" as just the duration of the next burst

Scheduling Algorithm Performance Criteria

- CPU Utilization - what % of time is the CPU executing a process?
- Wait time - how long (total) does a process spend in the ready queue?
- Turnaround time - time between arrival and completion?
- Response time - time between arrival and first useful response
- Throughput - number of processes completed per time unit.

First Come, First Served (FCFS)

- Jobs are executed in the order they arrived.
- Average waiting / turnaround times are order-sensitive.
- Pros:
 - Simple algorithm
 - Minimal context switching
- Cons:
 - Average waiting time typically poor
 - " " " is highly variable
 - CPU-bound processes can hog the CPU

Shortest Job First (SJF)

- When a scheduling decision is made, the shortest job is chosen.
 - NB - Jobs can only be selected once they join the ready queue.

Side Note: Pre-emptive Scheduling (PE)

- Preemptive: the current process can be stopped mid-execution and swapped for another
- Non-preemptive: once selected, a process is allowed to finish its CPU burst.

- Pre-emptive Shortest Job First is also known as:
Shortest Remaining Time First (SRTF)

- If a new job arrives that is shorter than the remaining time of the current job, it is executed.
- Pros:
 - long processes cannot hog the CPU (esp. with PE)
 - maximises throughput
 - average wait time is reduced
- Cons:
 - long processes can be starved
 - PE requires multiple context switches
 - execution time must be reliably estimated
 - overhead for maintaining a sorted queue.

Priority Scheduling

- Each job gets a priority number
 - Generally lower number = higher priority.
- Can be PE or non-PE.
- SRTF and SJF are special cases of this where the job time (or remaining time) is the priority.
- Pros:
 - high priority jobs (like user interactive programs) have short wait times
 - users/admins can have some control over scheduling
 - deadlines can be met.
- Cons:
 - risk of starvation for low-priority processes
 - can alleviate with an ageing system
 - more context switching if PE.
 - overheads for running a sorted queue.

Round Robin Scheduling (RR)

- Features a time quantum, q , usually around 10-100ms.
- When a process has run for this time it is pre-empted and rejoins the empty queue.
- Performance varies according to q :
 - If q is too high, RR tends towards FCFS
 - If q is too low, context switching overhead becomes significant.
 - q should be longer than about 80% of CPU bursts.
- Pros:
 - no starvation
 - predictable wait time
 - fair CPU sharing
- Cons:
 - poor average response time
 - waiting time linked to queue size, not priority
 - lots of context switching
 - hard to meet deadlines.

Multi-Level Queue Scheduling (MLQ)

- If processes can be grouped, we can use MLQ, e.g:
 - a queue for each priority level
 - two queues: foreground / background
 - etc
- The ready queue is split into multiple queues:
 - each process joins a queue
 - each queue can have a different scheduling algo.
 - an algo is needed to choose between queues.
 - often PE fixed priority

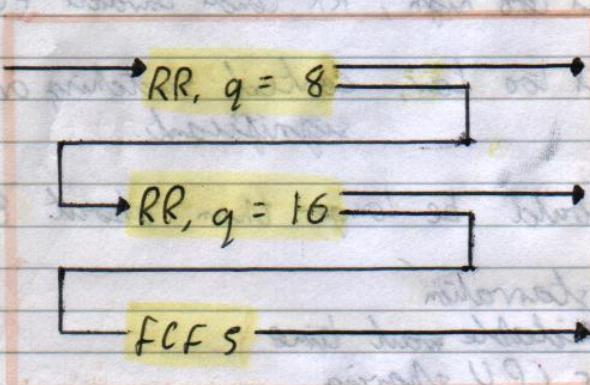
Multilevel Feedback Queue Scheduling (MLFQ)

- Like MLQ, but processes can move between queues
 - Used to implement ageing in priority scheduling
- We also require:
 - algo to decide when to promote/demote processes
 - algo to decide which queue is joined first by a process

Priority Scheduling w/ MLFQ:

- one queue for each priority level
 - Higher priority \Rightarrow more CPU time
- each process enters its correct priority queue
- after some time, if not complete, it is promoted.

Example:



Multi-Processor Scheduling

- Multiple CPUs can share the load
 - We assume all CPUs have the same capabilities
- Two variants: asymmetric and symmetric
- **Asymmetric Multiprocessing (ASMP)**
 - One processor makes all scheduling decisions, handles I/O processing and other system events.
 - Other CPUs execute user code only
 - The need for data sharing is reduced
- **Symmetric Multiprocessing (SMP)**
 - Each proc. does its own scheduling
 - May be one ready queue, or one for each processor
 - Each processor selects a process to execute from the ready queue.
 - High CPU-utilization requires **load balancing**:
 - **Push migration**: a specific process checks loads and redistributes processes as necessary
 - **Pull migration**: an idle processor pulls jobs from the queue of a busy processor.
 - All major OSs support SMP.

Hypertreading

- Multiprocessing requires multiple physical CPUs
- **Hypertreading** (aka. Symmetric Multithreading) uses multiple logical processors
 - Logical Processors:
 - Share CPU H/W: cache, bus, etc.
 - manage their own interrupt handling
 - Appears as two CPUs: one for floating point (FPU) and one for arithmetic and logic (ALU).

Scheduling Algo Evaluation

Deterministic modelling

- Run a set of processes through each algo.

Queueing Model

- Use queueing theory to predict utilization

- Little's formula: $n = \lambda \times w$

o n = queue length

o λ = arrival rate (per time unit)

o w = waiting time (time units)

- DO NOT USE unless given an arrival rate

Problems with these approaches:

- unrealistic; they don't use real processes
- hard to model complex algorithms with queueing theory
- arrival rate / waiting time hard to predict

Evaluation vs. Implementation

- The only accurate way to evaluate is to implement

o lots of work

o Users get inconsistent performance

o Processes change a lot

o Users will adapt to cheat.

- The ideal is to have a system where the admin can select the scheduling algo, specific to needs.

- An alternative is an API for users to change thread priorities.

Operations on Processes: Process Termination

- Processes can ask the OS to terminate it via exit

o Exit status value from a child is read by parent via wait()

o Process resources are deallocated by the system

- Parent may terminate a child using kill()

- If parent is exiting:

o Some OSs do not allow a child to continue

o All children are terminated - cascading termination.

Threading

- A thread is a **lightweight process**: a basic unit of CPU utilisation.
- It has a thread ID, program counter, register set and stack.
- It **shares code, data and OS resources** (like files) with other threads in the same process.
- If a process has multiple threads, it can do multiple tasks at once.

Process, kernel Threads and User Threads

Process

- Isolated with own virtual address space
- Contains process data like file handles
- **Lots** of overhead.
- Every process has at least one kernel thread

kernel Threads

- Shared virtual address space
- Contains running state data
- **Less** overhead
- From an OS PoV, this is what's scheduled for CPU time

User Threads

- Shared virtual address space
- Contains running state data
- **Even less** overhead
- Kernel unaware: appears as a **single threaded process**

- Main difference: **user threads** are controlled by application

Multi-Threading

- Most applications are multi-threaded

- This is good: other threads can continue while one is **blocked**.

- It also allows a **similar task** to be done **many times in parallel**, such as a server responding to requests.

- Technically most of this can be done with processes, but they have a **higher overhead**.
- **Multi-Threading Benefits:**
 - Responsiveness, as mentioned earlier
 - Resource sharing (and easy interaction)
 - Efficiency: threads are cheap to create and context switch.
 - Multi-processor utilization: a single-thread process only ever runs on one CPU core.

User Threads

- We could write code to appear as a single kernel thread that **internally schedules** many threads
 - Each "internal" thread is a **user thread**
- Alternatively, each "internal" thread can have its own **kernel thread**, allowing the OS to do the scheduling.

Threads in Java

- To **make a thread**, either **extend Thread** or **implement Runnable**.
 - Both have a "**run()**" method that must be implemented.
- To **start a thread**, either:
 - (**new ThingThatExtendsThread().start();**)
 - or
 - (**new Thread(new ThingThatImplementsRunnable()).start();**)
- Do not call **run()** directly!

Wasting Time

```
try {  
    Thread.sleep(30); // milliseconds  
} catch (InterruptedException e) {  
    e.printStackTrace();  
}
```

- Why? Prevents CPU hogging and allows another thread to take over.

- Thread.sleep() causes a blocking state, so a context switch happens.

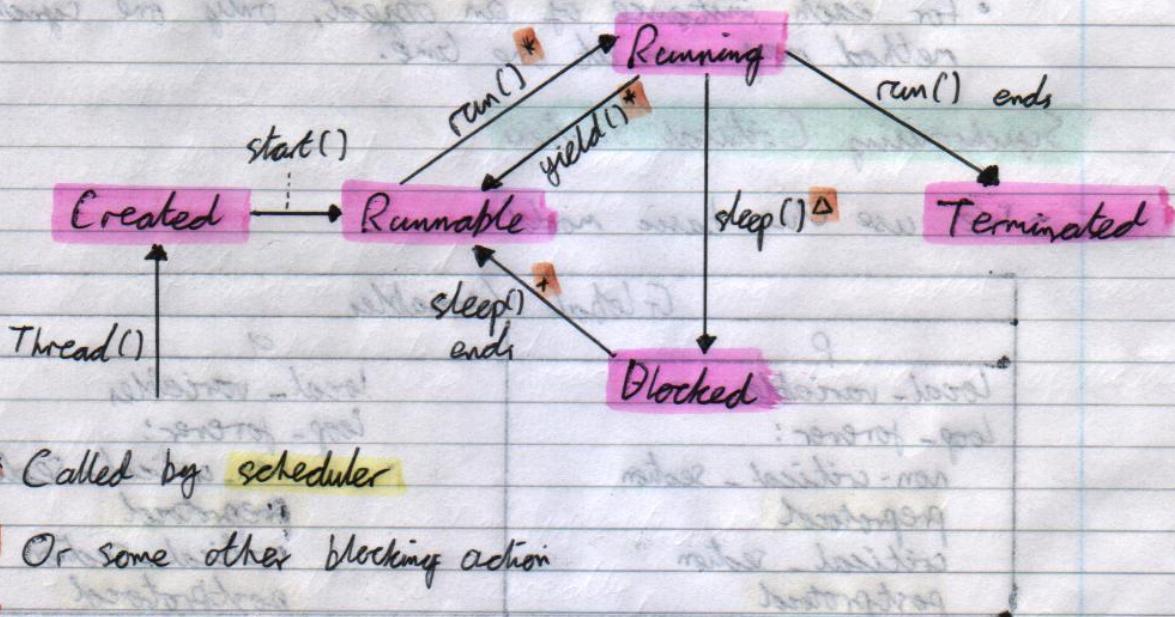
Waiting for Another Thread

q.join() → "Wait for q to finish before carrying on"

Killing Threads

- Asynchronous Cancellation: the thread is cancelled immediately.
- Deferred Cancellation: the thread checks periodically to see if it has been cancelled.
- A thread may be killed while it is sleeping, so the InterruptedException is thrown to allow for any cleanup.

Thread Lifecycle



Interleaving & Atomicity

- Threads share resources
- If two threads interact with the same resource, the interleaving of those threads means that conflicts may happen.
- Interleaving happens at runtime and is managed by the scheduler, so may be different for each run.
- An atomic statement is one that cannot be split/interrupted.
- Concurrency is the interleaving of atomic statements.
- We assume that each line of pseudo code is atomic (it isn't really!)

Critical Section

- A critical section must be executed atomically, wrt other related threads.
- Java provides the keyword `synchronized` to mark a method as a critical section.
 - For each instance of an object, only one synchronized method may run at one time.

Synchronising Critical Sections

- We use this basic model:

Global Variables	
p local-variables loop-forever: non-critical-section preprotocol critical-section postprotocol	q local-variables loop-forever: non-critical-section preprotocol critical-section postprotocol

preprotocol: check it is okay to enter CS }
postprotocol: signal that CS is done } synchronization mechanism

- Solution Properties

- Mutual exclusion: only one thread can enter a CS at one time
- No deadlock: if some threads are waiting to enter their CS, eventually one must succeed.
- No starvation: if a thread is waiting to enter its CS, it must be able to do so eventually.
- A good solution has no interleaving that breaks these rules.
- Standard testing is no good because we have to check every interleaving possible.

- Model Assumptions

- A thread will complete its CS.
 - No termination
 - No infinite loop
- A process may terminate during the non-CS
- Assignment is atomic
- CPU scheduler will not starve a process

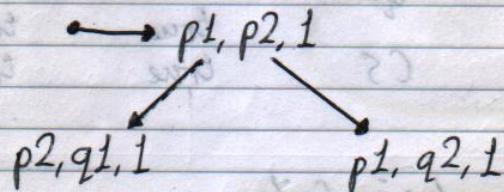
- First Attempt

int turn $\leftarrow 1$

P
loop:
p1: non-CS
p2: await turn = 1
p3: CS
p4: turn $\leftarrow 2$

Q
loop:
q1: non-CS
q2: await turn = 2
q3: CS
q4: turn $\leftarrow 1$

- State diagram:

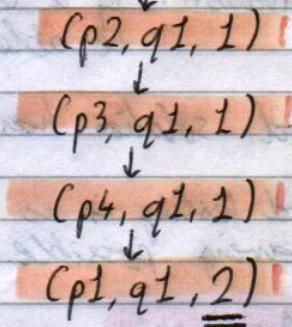


etc... The real tree is huge.

• Mutual Exclusion: there is no state $(p_3, q_3, -)$, so this holds!

• No Deadlock: deadlock can only happen at $(p_2, p_2, -)$, but at those states either turn=1 or turn=2, so eventually a thread proceeds.

• No Starvation: Suppose we have: $(p_1, q_1, 1)$



then q quite during its non-CS. p will never execute now!

- Second Attempt

	p	$wantp \leftarrow \text{false}$	$wantq \leftarrow \text{false}$	q	$wantp \leftarrow \text{false}$	$wantq \leftarrow \text{false}$
loop:				loop:		
non-CS				non-CS		
p_1 : await $wantq = \text{false}$				q_1 : await $wantp = \text{false}$		
p_2 : $wantp \leftarrow \text{true}$				q_2 : $wantq \leftarrow \text{true}$		
CS				CS		
p_3 : $wantp \leftarrow \text{false}$				q_3 : $wantq \leftarrow \text{false}$		

• Possible Interleaving:

p	q	$wantp$	$wantq$
p_1		false	false
	q_1	false	false
p_2		true	false
	q_2	true	true
CS		true	true
	CS	true	true

Mutual exclusion fails.

- Third Attempt

Third Attempt -

\rightarrow wantp \leftarrow false now wantq \leftarrow false

P
loop:
p1: non-CS
p2: wantp \leftarrow true
p3: await wantq = false
p4: CS
p5: wantp \leftarrow false

q
loop:
q1: non-CS
q2: wantq \leftarrow true
q3: await wantp = false
q4: CS
q5: wantq \leftarrow false

• Mutual Exclusion (Informal Proof)

- We must show that (p_4, q_4, \dots, \dots) cannot occur.
- To get that, one thread must be at 4 and the other must move $3 \rightarrow 4$.
 - Threads are symmetric, so it doesn't matter which one.
- Let's say p is at p_4 . What is the value of wantp?
 - q cannot modify wantp.
 - If p is at p_4 , the last change was p2: wantp \leftarrow true
- For q to move $q_3 \rightarrow q_4$, we need wantp = false
 - But p4 only allows wantp = true.
- ∴ (p_4, q_4, \dots, \dots) is not possible!

• No Deadlock

P	q	wantp	wantq
p1	q1	false	false
p2	q2	false	false
p3	q3	true	true
		true	true
		true	true

Deadlock!

- Peterson's Algorithm

wantp \leftarrow false, wantq \leftarrow false, last $\leftarrow 1$

P

- p1: non-CS
- p2: wantp \leftarrow true
- p3: last $\leftarrow 1$
- p4: await wantq = false or
last = 2
- p5: CS
- p6: wantp \leftarrow false

q

- q1: non-CS
- q2: wantq \leftarrow true
- q3: last $\leftarrow 2$
- q4: await wantp = false or
last = 1
- q5: CS
- q6: wantq \leftarrow false

• Mutual Exclusion (Informal Proof)

- We must show that $(p_5, q_5, -, -, -)$ cannot occur.

- We must have one process at S and one moving 4 \rightarrow 5 for that to happen.

- Assume p is at p5: wantp = true (see p2)

- We know $(p_5, q_4, \text{true}, -, 1 \text{ or } 2)$

- We want to show that $q_4 \rightarrow q_5$ cannot happen.

- wantp = true, so $q_4 \rightarrow q_5$ can only happen if last = 1, which can only be set at p3.

- When $p_4 \rightarrow p_5$ happened, either:

o wantq = false

- so q was at q1 or q2

- and q must exec. q3 before q4, setting last = 2

- p is at p5 so cannot set last = 1 at p3

- so q waits at q4

o last = 2

- p is at p5 so cannot set last = 1 at p3

- so q waits at q4

◦ No Deadlock (Informal Proof)

- Both must be stuck at line 4.
- $want_p$ and $want_q$ are true, and don't change.
- Because $last$ must = 1 or = 2, the program cannot get stuck here.

◦ No Starvation (Informal Proof)

- We must show that if p is at p_4 it will eventually reach p_5 .
- If p is at p_4 , it can proceed unless $want_q = \text{true}$ and $last = 1$.
- Where could q be?
 - If $want_q$ is true, q can be at q_3 , q_4 or q_5
 - If q is at q_3
 - it will execute $last = 2$
 - it will remain at q_4 until p executes
 - If q is at q_4
 - it will move to q_5 because $last = 1$
 - it will proceed to q_6
 - it will proceed to q_1
 - it may quit here; p can execute because $want_q = \text{false}$
 - it will proceed to q_2 , setting $want_q = \text{true}$
 - it will proceed to q_3 , setting $last = 2$
 - it will remain at q_4 until p executes.
 - If q is at q_5

Properties as Temporal Logic

◦ Mutual exclusion: $\square \neg(p_4 \wedge p_5)$

◦ No Deadlock: $\square ((p_4 \vee q_4) \rightarrow \diamond (p_5 \vee q_5))$

◦ No Starvation: $\square (p_4 \rightarrow \diamond p_5)$

Concurrency with n Processes

Baking Algorithm

- Based on a service in a bakery or other shop.
- Each process takes a number; **lowest number goes first**.
- For N processes:

```
int array [1...N] ← [0, ..., 0]
```

loop forever:

p1: non-CS

p2: number [i] ← max(number) + 1 *

p3: for all other processes j :

 await number [j] = 0 or number [j] < number [i]

p4: CS

p5: number [i] = 0

* This is not atomic! This can be fixed, but the algo is more complex.

Good:

- each variable is written to by only one process
- satisfies all three properties

Bad:

- unbounded ticket numbers
- p3 queries every other process.

Ugly:

- too inefficient to actually use.

Synchronization Hardware

- So far, we have assumed an atomic assignment operator.

- Now consider a more general idea:

"A thread must obtain a lock to enter its CS"

- Hardware can provide these locks as primitives to make synchronization easier.

Side Note: in a **single processor system** we could solve the CS problem by **disabling interrupts** during a CS, but in a **multi-proc system** this message has to be passed to every core, which is slow.

Test and Set

boolean lock \leftarrow false

loop - forever:

p1: non-CS

p2: await ($! \text{testAndSet}(\text{lock})$)

p3: CS

p4: lock \leftarrow false

boolean testAndSet(boolean l) {
 boolean toReturn =
 l.getValue();
 l.setValue(true);
 return toReturn;
}

- Implemented as a single processor instruction
- Does no harm on set if already true
- False \rightarrow no one is in CS

Swap

boolean lock \leftarrow false

P

boolean key \leftarrow true

loop - forever:

p1: non-CS

p2: key \leftarrow true

p3: while ($\text{key} = \text{true}$):
 swap (lock, key)

p4: CS

p5: lock \leftarrow false

boolean swap(boolean a, boolean b) {
 boolean tmp = a.getValue();
 a.setValue(b.getValue());
 b.setValue(tmp.getValue());
}

I implemented as a single processor instruction.

Both of these could lead to starvation, but it is unlikely.

They solve deadlock and mutual exclusion.

Semaphores

- Programmers usually prefer high-level, software mechanisms for concurrency control:
 - o More abstract
 - o Less complex
 - o Less error prone.
- Semaphores are provided by most OSs.

Busy-Wait Semaphore

- Contains:
 - o An integer variable, v (always $v \geq 0$)
 - o A set of blocked processes (initially empty)

- Has two methods:

```
wait (Semaphore s) {
    while (s.v <= 0) {
        // do nothing
    }
    s.v = s.v - 1;
}
```

```
signal (Semaphore s) {
    s.v = s.v + 1;
}
```

* also a "spin-lock"

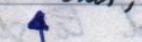
Blocked-Set Semaphore

- Contains the same as above

- Has two methods:

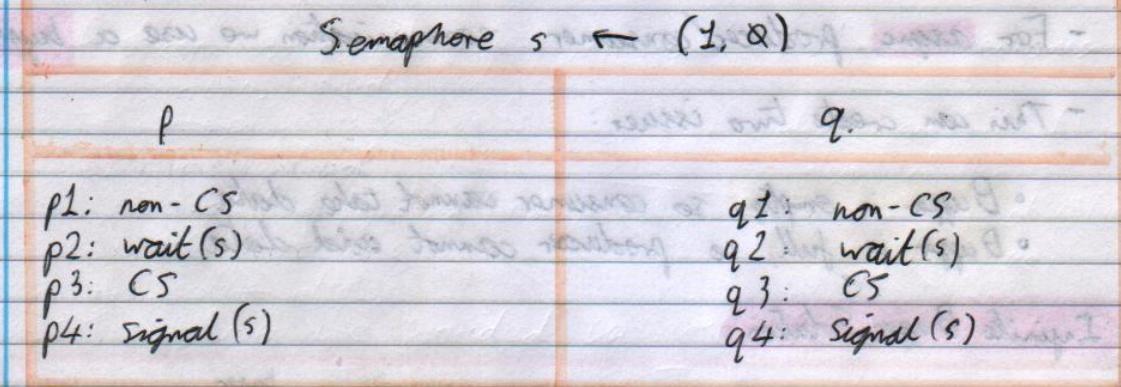
```
wait (Semaphore s) {
    s.v = s.v - 1;
    if (s.v < 0) {
        s.blocked = s.blocked ∪ this;
        this.state = blocked;
    }
}
```

```
signal (Semaphore s) {
    s.v = s.v + 1;
    remove one q. from s.blocked;
    q.state = runnable;
}
```



Union, i.e. add
this to s.blocked.

Semaphore Usage



- Initially, $s.v = 1$
- Before a thread enters CS:
 - It calls wait; $s.v$ is decremented; 0 means some thread in CS
 - If another thread calls wait, then $s.v = 0$ so it is blocked.
- When a process leaves CS:
 - signal(s) adds one to $s.v$, so another process can enter
 - if there are waiting processes, one is unblocked.

Properties & Notes

- Properties:

- In busy-wait, $s.v > 0$ always.

- $s.v = \text{init} + \# \text{signal}(s) - \# \text{wait}(s)$

- $\#CS + s.v = 1$

- $\#CS = \# \text{wait}(s) - \# \text{signal}(s)$

- Mutual Exclusion ✓

- Freedom from Deadlock ✓

- Starvation: technically possible, but requiring a specific interleaving so unlikely. Can be fixed with a busy FIFO queue, instead of a blocked set.

- Creating a semaphore with a value higher than 1 allows a control such as "max of 3 CSs at once".
 - Used for server connections.

The Producer/Consumer Problem.

- For **asynchronous** producer/consumer communication we use a **buffer**.
- This can create two issues:
 - Buffer is empty, so consumer cannot take data
 - Buffer is full, so producer cannot add data

Infinite Buffer Solution

Semaphore notEmpty $\leftarrow (0, \emptyset\right)$, Buffer $\langle d \rangle^{\text{buffer}} \leftarrow \emptyset$	
producer	consumer
p1: $d \leftarrow \text{produce}()$ p2: $\text{append}(d, \text{buffer})$ p3: $\text{signal}(\text{notEmpty})$	c1: $\text{wait}(\text{notEmpty})$ c2: $d \leftarrow \text{take}(\text{buffer})$ c3: $\text{consume}(d)$

- We only need to control removal
- Semaphore starts at zero because we want the consumer to wait until something has been added
- $\text{Semaphore} = 0 \Rightarrow$ buffer is empty.

Finite Buffer Solution

Semaphore notEmpty $\leftarrow (0, \emptyset\right)$, Semaphore notFull $\leftarrow (N, \emptyset\right)$, size N queue $\langle d \rangle^{\text{buffer}} \leftarrow \emptyset$	
producer	consumer
p1: $d \leftarrow \text{produce}()$ p2: $\text{wait}(\text{notFull})$ p3: $\text{append}(d, \text{buffer})$ p4: $\text{signal}(\text{notEmpty})$	c1: $\text{wait}(\text{notEmpty})$ c2: $d \leftarrow \text{take}(\text{buffer})$ c3: $\text{signal}(\text{notFull})$ c4: $\text{consume}(d)$

- notFull is initialised to N , so N things can be added before it starts blocking.
- This technique is called **split semaphores**.

Semaphores in Java

Contract: new Semaphore (int permits); → fair = false.

new Semaphore (int permits, boolean fair);

Wait:

s. acquire();

Signal:

s. release();

Monitors + Concurrency in Java

- Semaphores are great (v. user-friendly), but prone to errors.
- Semaphores must be implemented perfectly and carefully, otherwise big errors can occur.

- Solution: Monitors

- Similar to a class, a monitor encapsulates a resource and provides access only via public methods.
- Responsibility for mutual exclusion is contained within the monitor for a resource.
 - Less scope for bugs
 - Less duplicated code
- Each resource can be handled by a separate monitor.

Implementing Monitors w/ Semaphores

```
class Monitor {
```

```
    Semaphore s = new Semaphore(1);
```

```
    method m1() {
```

```
        s.wait();
```

```
        // actual method code
```

```
        s.signal();
```

```
}
```

```
    ...
```

```
}
```

Simulating Semaphores w/ Monitors

monitor Semaphore {

 int s = k;

 method semaWait()

 while (s == 0)

 wait();

 s = s - 1;

 method semaSignal()

 s = s + 1

 notifyAll();

}

P

Q

loop-forever:

p0: non-CS

p1: Semaphore.semaWait()

p2: CS

p3: Semaphore.semaSignal()

loop-forever:

q0: non-CS

q1: Semaphore.semaWait()

q2: CS

q3: Semaphore.semaSignal()

Monitor wait() and notifyAll()

- Maintain a list of waiting blocked threads

- **wait():**

- adds the current thread to the blocked list
- sets its state to blocked
- releases the monitor lock

- **notifyAll():**

- sets all blocked states to ready
- empties the blocked list.

→ the scheduler will select one - the rest will re-block.

- These can only be called **inside** the monitor.

Producer/Consumer Problem w/ Monitors

monitor Buffer

int [5] - buffer;

int spaceUsed = 0;

method addItem (i)
while (spaceUsed == 5)
 wait();
 buffer[spaceUsed] = i;
 ++ spaceUsed;
 notifyAll();

method removeItem ()
while (spaceUsed == 0)
 wait();
 i = buffer[spaceUsed];
 -- spaceUsed;
 notifyAll();
 return i;

P

loop - forever:
p0: i ← produce();
p1: Buffer.addItem(i);

q

loop - forever:
q0: i ← Buffer.removeItem();
q1: consume(i);

- It would be better if we could **notify threads selectively**.

◦ i.e. only wake producers when the buffer is empty.

- Classical monitors allow this, but Java does not.

◦ **wait()** takes a condition that the thread is waiting for
◦ **notify/notifyAll()** takes a condition to signal

Queue or Set in Monitors

- Only worthwhile if we'll use **notify()** to wake the head of the queue.

- This can avoid starvation caused by allowing the scheduler to select a thread after **notifyAll()**;

- Need to be careful though: this only works if all threads are waiting for exactly the same condition.

Can Monitors Deadlock?

- Yes.

- But the deadlocking code is all in one place - easier to debug!

Notifying Before Returning

- We can't always make `notifyAll()` the last command
• What if we have to return?

- What happens if the thread is in the monitor and notifies other threads?

Thread States w.r.t Monitors:

- E = waiting to enter the monitor
- W = waiting to be notified
- N = executing notify (currently holds the lock)

- In Java, priority is $E = W < N$

- Traditionally, it is $E < N < W$

- $E < N < W$ Implications

• If a single process is woken then it knows the condition it was waiting for is true because it has just been signalled: no thread has executed between it and the one that signalled it (which has now stopped) so it can continue without rechecking the condition.

• This is the immediate resumption requirement.

• $E = W < N$ Implications

• Good if we want to return before we notify: the signalling process can continue until it has left the monitor.

• Now there is no guarantee that the signalling process didn't change the wait condition before it stopped.

• Further, `notifyAll()` could have caused a different process to wake and break the condition.

- In this case, each thread must re-check the condition before continuing

- Achieved by using `while (! condition)` - it cannot exit until the condition has been checked and seen to be true.

Monitors in Java

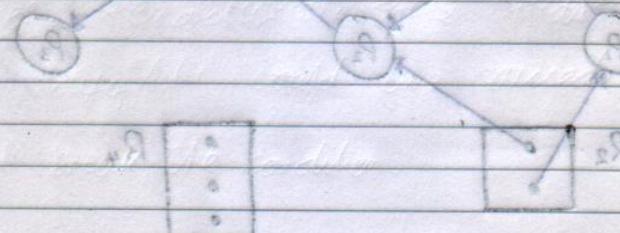
- Don't really exist.
- All Objects can be used as a monitor
- Not all methods must be mutually exclusive, but those that require ME are marked with `synchronized`.
- All Objects implement `wait()`, `notify()` and `notifyAll()`
- Refer to start of notes for Java thread lifecycle.

Deadlock

- Deadlock occurs during concurrent programming when, for some reason, no process is able to continue.
- It occurs in interaction between threads/processes, so the cause may not be in one piece of the code.

Conditions for Deadlock

- System Model :
 - A number of processes.
 - A number of resources
 - Disk
 - Semaphores
 - etc.
 - Each process requests resources as it needs them
 - ... " releases " " finishes "
 - Some processes use multiple resources with them
 - Some resources have multiple instances



- Cogman Conditions for Deadlock:

- Mutual exclusion - only one process at a time can use an instance of a resource
- Hold and Wait - a process holding at least one resource is waiting to acquire additional resources held by other processes.
- No Preemption - a resource is only released voluntarily by the process holding it, after it has completed its task.
- Circular Wait - there exists a set of processes $\{P_0 \dots P_n\}$ such that:
 - P_0 is waiting for a resource held by P_1
 - P_1 " " " " " held by P_2
 - P_2 " " " " " held by P_3
 - P_{n-1} " " " " " held by P_0
 - P_n " " " " " held by P_0

Resource Allocation Graphs

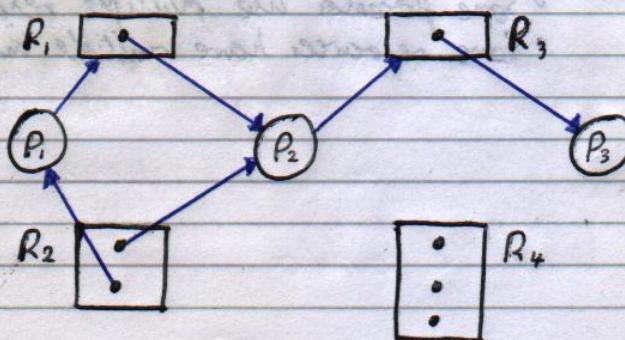
- Vertices: processes (○) and resources (□)
- Request Edge: $(P_i) \rightarrow R_j$ = P_i has requested an instance of R_j
- Assignment Edge: $(R_i \circ \circ) \rightarrow (P_j)$ = P_j holds an instance of R_i
- Example:

$$P = \{P_1, P_2, P_3\}$$

$$R = \{R_1, R_2, R_3, R_4\}$$

Num. of instances: $R_1: 1, R_2: 2, R_3: 1, R_4: 3$

$$E = \{P_1 \rightarrow R_1, P_2 \rightarrow R_3, R_1 \rightarrow P_2, R_2 \rightarrow P_2, R_2 \rightarrow P_1, R_3 \rightarrow P_3\}$$



- RAG Cycles

processes to wait forever waiting when

- No cycles = no deadlock

◦ If cycles exist...

- If each resource has only one instance, we have deadlock

- If some resources have multiple instances, we may have deadlock.

Approaches to Deadlock

- Prevention: ensure the system cannot enter a deadlocked state

- Cure: allow deadlock to happen and then recover

- Ignore it!

◦ Dealing with it is expensive, and it's pretty rare anyway.

Preventing Deadlock

- We must break one of the Coffman conditions:

- Mutual exclusion: not easy, but we can reduce the risk of a problem by enforcing ME when only absolutely necessary.

- Hold and wait: guarantees that whenever a process requests a resource it does not hold any other

◦ a process must request and hold all resources before it can begin

◦ or a process can only request a resource when it has none

◦ causes low resource utilization and can cause starvation.

- Pre-emption: - if a process requests a resource that cannot be allocated immediately it must immediately release any resources already held

◦ a list of all resources the process is waiting for is kept

◦ the process is only restarted when all resources it needs are available.

- **Circular Wait** - Impose a total ordering on resources and make processes request them in increasing order.
- To avoid a deadlock we need extra information about processes:
 - e.g. max number of resources of each type that will be requested
 - Unrealistic to assume will know this.
- We can use the resource allocation state of a system to ensure that we will never allow a circular wait
 - With single-instance resources we can do this with a RAG.
 - With multi-instance resources we use Banker's algorithm.
- **Resource Allocation State:** number of available/allocated resources and max number of resources each process may need.

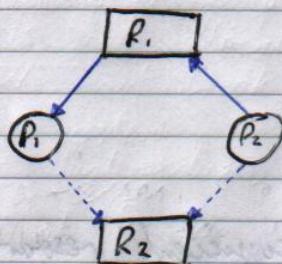
Safe States

- Each time a process requests a resource, the system must decide if that allocation leaves the system in a **safe state**.
- A system is in a safe state if there exists a sequence of all processes $\langle P_0, P_1, \dots, P_n \rangle$ such that for all P_i , the resources that P_i can request can be satisfied by currently available resources and the resources held by all P_j with $j < i$.
- Therefore...
 - A process can never be holding a resource and waiting for a process with a higher number to release one.
 - Circular wait cannot occur: P_j can hold and wait for P_i , but P_i cannot hold and wait for P_j . ($j > i$)
- Safe state = cannot deadlock
- Unsafe state = might deadlock

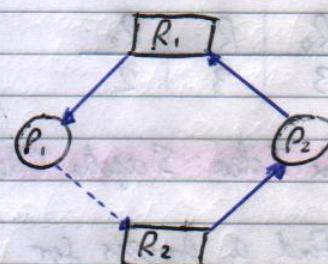
Safe States - RAG Method

- Works only with single-instance resources
- Add a new edge type: a claim edge. (dashed line)
- A claim edge from $P_i \rightarrow R_j$ shows that P_i can request R_j . This is specified a priori.
- Claim edge changes to assignment edge and back when R_j is allocated/released.
- All claim edges must be defined a priori.
- If a process P_i requests resource R_j , the request can only be granted if the converting the $P_i \rightarrow R_j$ claim edge to an assignment edge does not create a cycle.

Safe:



Unsafe:



Safe States - Banker's Algorithm

- Works with multi-instance resources.
- Assumption: resources are eventually released.
- Processes may have to wait after they make a request.
- Not cheap: $O(mn^2)$ for m resources and n processes.

- Input Data Structures

◦ For n processes and m resources

◦ Allocation : $n \times m$ matrix

Allocation $[i, j] = k$ means

P_i has k instances of

R_j

◦ Max : $n \times m$ matrix

Max $[i, j] = k$ means P_i may

request upto k instances of R_j

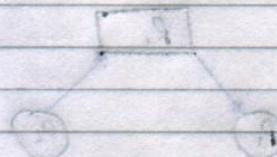
	R_0	R_1	R_2
P_0	0	1	0
P_1	2	0	0
P_2	3	0	2
P_3	2	1	1
P_4	0	0	2

	R_0	R_1	R_2
P_0	7	5	3
P_1	3	2	2
P_2	9	0	2
P_3	2	2	2
P_4	4	3	3

◦ Available : m -length vector

Available $[i] = k$ means there are k instances of R_i available now

	R_0	R_1	R_2
	3	3	2



- Basic Idea: Safety Check

1. Find a process for which there are sufficient resources to meet its max need.

2. Simulate executing the process by releasing all of its resources.

3. Go back to step 1. Collect \$200.

4. If all processes have completed, the system is safe.

5. If any processes remain, the system is unsafe.

- Algorithmic Safety Check

1. Let $work$ and $finish$ be arrays of m and n respectively:

- $work = available$

- $finish[i] = false$ for $i = 0, 1, \dots, n-1$

→ Resources
→ Processes

2. Find any i such that:

- $finish[i] = false$ and $\text{max}[i] - \text{allocation}[i] \leq work$

- If no i exists, go to step 4.

3. $\text{Work} = \text{Work} + \text{allocation}[i]$
 $\text{finish}[i] = \text{true}$
Go back to step 2.

4. If all entries in finish are true , the system is **safe**.
Otherwise, it is **unsafe**.

- How to Use This?

- When a process P_i requests a resource R_j :
 - Update allocated as if the request was granted
 - Run the algo.
 - If the system is **safe**, allow the request. Otherwise, P_i must **wait**.

Deadlock Detection

- Instead of prevention, we can **allow it to happen** and then **recover**.
- With **single-instance resources** we can look for cycles in the **RAG**.
- With **multi-instance resources** we can use a modified **Banker's algo**.
 - Instead of using the **max number of resources** a process can request, use the **number it is currently requesting**.
- Banker's Algo for Deadlock Detection**

1. Initialize as before, but if P_i has no resources allocated then $\text{finish}[i] = \text{true}$.

2. Find any i such that:

- $\text{finish}[i] = \text{false}$ and $\text{request}[i] \leq \text{work}$
- If no such i exists, go to step 4.

3. $\text{work} = \text{work} + \text{allocation}[i]$
 $\text{finish}[i] = \text{true}$
Go to step 2.

4. If $\text{finish}[i] = \text{true}$ for all i , the system is **safe**.
If $\text{finish}[i] = \text{false}$ then process P_i is **deadlocked**.

- In step 2 and 3 we assume that if P_i has enough resources to continue now then it will complete. It may actually request more resources a little later, but that's okay: we will detect a deadlock then, if it happens.

- When to Call This?

- When, and how often, depends on:
 - How often a deadlock is likely to occur?
 - How many processes will need to be rolled back?
 - A: one for each disjoint cycle.
- If not frequent enough then ~~many~~ many cycles could be created and we wouldn't know which was the "cause".
- If too frequent then it is a waste of CPU time.
- Assuming we detect a deadlock, how do we recover?

Deadlock Recovery

- Process termination

- Abort all deadlocked processes
- Abort one at a time until
 - What order to abort in?
 - Process priority:
 - Time elapsed / remaining
 - Resources used
 - Resources required
 - Process interactive or batch?

- Resource Pre-emption

- Select a process and take resources away to break the deadlock
- Rollback to a safe state and restart the process(es).
- Select a victim that minimizes cost.
- Include number of rollbacks in cost factor to prevent starvation by always selecting the same process.