

# DSM: Tutorial 5

## Question 1

Improvement options:

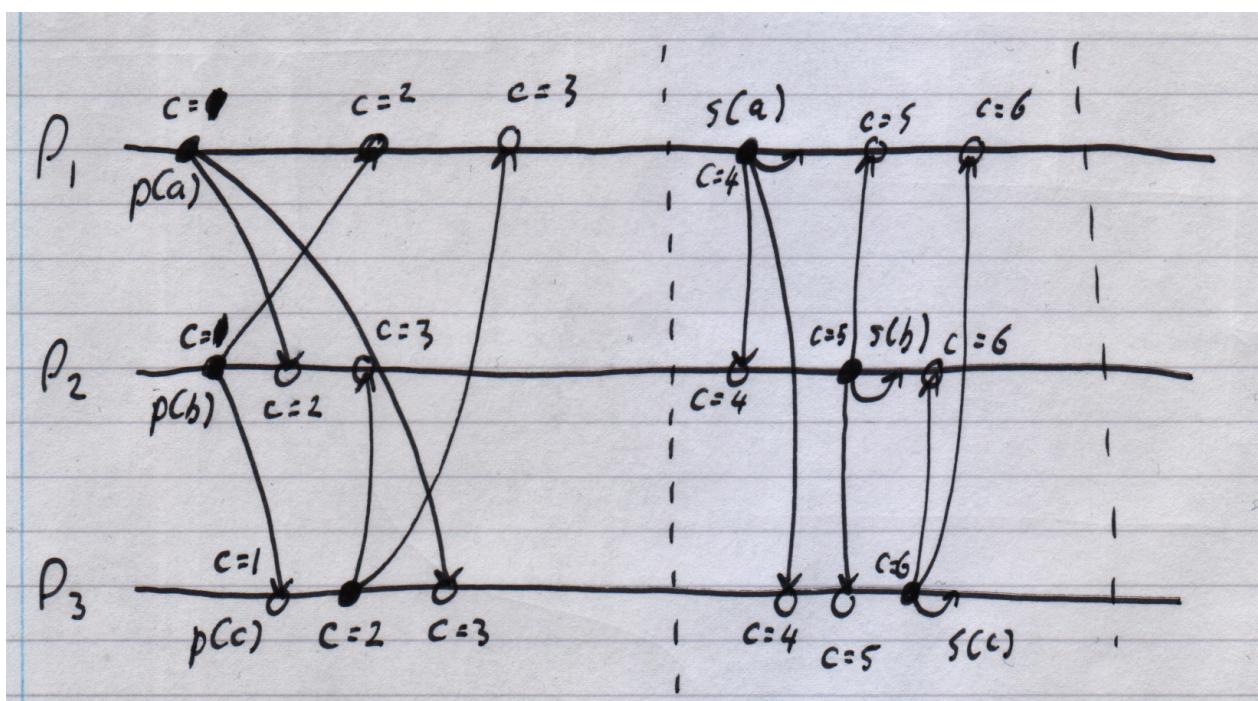
- Change the consistency ordering to reduce the blocking time, at the cost of lost requests.
- Don't bother with acknowledgements to reduce the blocking time, at the cost of lost requests.
- Allow a front-end to talk to back-ups to allow for load balancing, at the cost of consistency guarantees being lost.

Options for recovering from inconsistency:

- Detect inconsistency with checksums.
- Periodically copy the primary (may overwrite data).
- Merge differing states to capture all requests (may lose data).

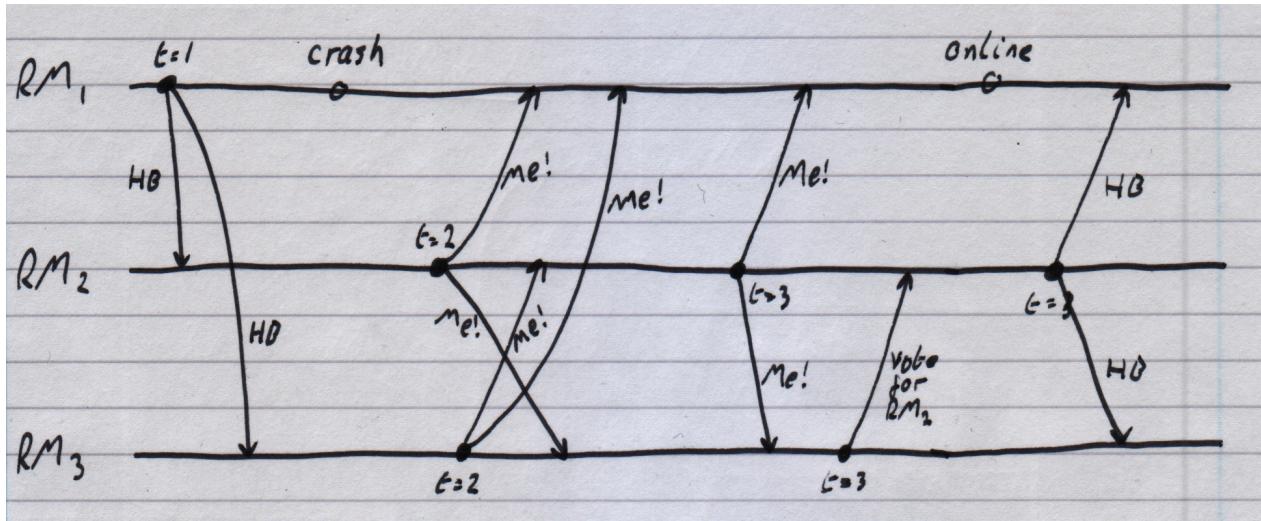
## Question 2

In the first stage, an arbitrary messaging protocol is used for each process to propose one of a, b or c. In the second stage, each process seconds a value chosen from the ones it received, this time with total ordering. Each process decides on the first value it received, so a is decided.



To avoid bias, we need to make sure that no process has an advantage. Total ordering uses processor IDs to break ties when multiple messages have the same clock, which creates a bias. This could be mitigated by changing processor IDs periodically (which would also require consensus, but it doesn't need to be unbaised if the proposals are random).

## Question 3



Any heartbeats from  $RM_1$  when it comes back online will be ignored. When it receives a heartbeat with a higher term, it will know it is not the leader.

## Question 4

A bug could introduce inconsistency at any point: incorrect/inconsistent parsing of input, failure to write data, etc. The inconsistency could be detected by comparing replicas (maybe using checksums) and overwriting one replica with data from another to achieve a consistent state (a process that may favour the newest copies of data).

## Question 5

| Primary Backup  | State Machine   |
|---|---|
| Requests are sent to one RM.  | Requests are sent to all RMs.   |
| When a backup RM fails, no pending requests should be lost. When the primary RM fails, an unbounded number will be lost during the failover period. | When one RM fails, no pending requests should be lost because other RMs will still receive the request. |
| For failstop errors, $t + 1$ RMs are needed for continually functioning service.  | For failstop errors, $t + 1$ RMs are needed for continually functioning service.                        |
| Byzantine errors cannot be handled.   | For Byzantine errors, $2t + 1$ RMs are needed for continually functioning service.                      |