

PAL: Parallel Algorithms

Important Notes About These Notes

These notes were written by me, Mark Ormesher, during my revision for Computer Science exams at King's College London. I am publishing them here to help my classmates and other students in the years below me. To the best of my knowledge they are correct.

- These notes are **not endorsed** by King's College London, the module lecturers, or any member of College staff.
- These notes are **not checked** by any qualified individual, and I offer no guarantee as to how accurate or complete these notes are. They are offered for free, as-is, and I accept no responsibility for how you choose to use them.
- These notes were relevant for my year (2016/17) but **the content for your course may be different**. Check your lecture slides, syllabi, tutorial guides, etc.
- These notes were produced for my own personal use (it's how I like to study and revise). That means that some annotations may be irrelevant to you, and some topics might be skipped entirely.
- Feel free to **share** these notes, however please only share a link to the repo (see the link below), not individual files.

Notes are originally from **<https://github.com/markormesher/CS-Notes>** and remain the copyright-protected work of Mark Ormesher.

List of PAL Algorithms

Multi-threading

- Loops
- Tree traversal
- Exponentiation
- Horner's method
- Merge sort (and fast merge)

PRAM EREW

- Fan-in array binary operation (sum, product, min, etc.)
- Broadcast
- Array membership (partial)
- Array membership (full)
- Array prefix sums

PRAM CR*W

- List ranking (CW or EW)
- Root finding (CW or EW)
- Matrix * vector multiplication (CW or EW)
- Matrix * matrix multiplication (CW or EW)
- Array sum (CW-priority)
- Array min/max (CW-any)
- Array membership (find index with CW-priority, find true/false with CW-arbitrary)
- Array sorting (CW must "stack")

Graph Interconnected Model

- 1D mesh even/odd transposition sort
- 2D mesh snake sort
- 2D mesh broadcast
- 2D mesh array min/max
- 2D mesh membership
- 2D mesh array prefix sums

Distributed Networks

- Directed ring leader election
- General leader election
- General breadth-first search tree finding

Graphs

- Sequential graph components
- Parallel graph components
- Parallel minimum spanning tree (Boruvka's method)

Multi-Threading Loops

Input: indexes i and j , $i \leq j$
some code to execute, exec .

Output: run exec for all indexes from i to j , inclusive

func $\text{paraLoop}(i, j, \text{exec})$:

if ($i == j$):
 $\text{exec}(i)$

} Base case: ran exec for the single index i (j would also work).

else:

$k = i + \text{floor}((j-i)/2)$ } Find the mid-point between i and j .

spawn $\text{paraLoop}(i, k, \text{exec})$ } Do half of the range on a new thread...

$\text{paraLoop}(k+1, j, \text{exec})$ } ... and half on this thread.

sync } Make sure everything finishes.

The runtime is $\Theta(\log_2 n)$ if n is the range from i to j .

Note: the overall runtime is $O(m \log_2 n)$ where m is the runtime of exec .

Multi-Threading Tree Traversal

Input: binary tree root node, t .

Output: pre-, in- or post-order traversal of the tree rooted at t .

func paraTraversal (t):

if ($t == \text{null}$): } Base case.
return []

$m = t.\text{data}$ } Visit this node.

$l = \text{spawn paraTraversal} (t.\text{left})$ } Traverse the left sub-tree on a new thread.

$r = \text{paraTraversal} (t.\text{right})$ } Traverse the right sub-tree on this thread.

sync } Wait for both traversals to finish.

For pre-order:
return $[m, l, r]$ } Implicit array flattening is assumed.

For in-order:
return $[l, m, r]$

For post-order:
return $[l, r, m]$

The span is $\Theta(h)$ where h is the height of the tree. For a balanced tree with n nodes, $h = \log_2 n$.

Multi-Threading Exponentiation

Input: integers n and k , $k \geq 1$

Output: n^k

func paraExp(n, k):

if ($k == 1$):
 return n

} Base case.

$j = \text{floor}(k/2)^*$

} We will compute two halves of
the problem separately.

$a = \text{spawn paraExp}(n, j)$

} The first half runs on a
new thread...

$b = \text{paraExp}(n, k-j)$

} ... and the second half runs on
this thread. Note " $k-j$ " to avoid
rounding errors from the point
marked *.

sync

} Wait for both halves to finish.

return $a * b$

} Re-combine and return the result.

The span is $\Theta(\log_2 k)$, assuming multiplication as an $O(1)$ operation.

Multi-Threading Horner's Method (Polynomial Evaluation)

Input: polynomial in array form, a Eg. $3x^2 + 4x + 9 = [9, 4, 3]$
value for x , x
indexes s and e

Output: value of the polynomial between the powers s and e inclusive.

Note: initially, $s=0$ and $e=\text{highest power in polynomial}$

func paraHorner(a, x, s, e):

if ($s == e$): } Base case.
 return a[s]

$m = s + \text{floor}((e-s)/2)$ } Find the midpoint of s and e.

a = spawn paraHorner(a, x, s, m) } Compute the range s to m
on a new thread.

b = paraHorner(a, x, m+1, e) } Compute the range m+1 to e
on this thread

sync } Wait for both halves to finish.

$\exp = m - s + 1$ } b will be too small by a factor
of x^{\exp} ...

return $a + (b * x^{\exp})^*$ } ... so we shift it up
accordingly.

The recursion depth and span are $\Theta(\log_2 n)$ for a polynomial with n terms. The overall runtime is $\Theta(\log_2^2 n)$ if exponentiation (marked *) can be done in logarithmic time with the parallel method.

Multi-Threading Array Merge.

Input: Two sorted arrays, a and b .

Output: The sorted merge of a and b .

func paraMerge(a, b):

$$m = \text{floor}((a.\text{length} + 1) / 2) \quad \} \text{Mid-position in } a.$$

$$x = a[m] \quad \} x \text{ is the pivot value.}$$

$$\begin{aligned} a_{\text{Low}} &= a[1 \dots m-1] \\ a_{\text{High}} &= a[m+1 \dots \text{end}] \end{aligned} \quad \} \text{Split } a \text{ into values above } m \text{ and values below } m.$$

- All values in a_{Low} are $\leq x$
- All values in a_{High} are $\geq x$

$$q = \text{findPositionWithBinarySearch}(b, x) \quad \} \text{Find the position of } x \text{ in } b, \text{ or its closest value above. below. above.}$$

$$\begin{aligned} b_{\text{Low}} &= b[1 \dots q-1] \\ b_{\text{High}} &= b[q \dots \text{end}] \end{aligned} \quad \} \begin{array}{l} \text{- All values in } b_{\text{Low}} \text{ are } \leq x \\ \text{- All values in } b_{\text{High}} \text{ are } \geq x \end{array}$$

$$l = \text{spawn ParaMerge}(a_{\text{Low}}, b_{\text{Low}}) \quad \} \text{Merge the Low end in one new thread.}$$

$$r = \text{paraMerge}(a_{\text{High}}, b_{\text{High}}) \quad \} \text{Merge the high end in this thread.}$$

sync } Wait for both merges to finish.

return [l, x, r] } Return the recombined merged set.
Note: implicit array flattening.

The recursion depth and span are $\Theta(\log_2 n)$, each level of which uses an $\Theta(\log_2 n)$ binary search.

The overall runtime is $\Theta(\log^2 n)$.

Multi-Threading Array Merge Sort

Input: array, a
start and end indexes, s and e

Output: a, sorted between s and e inclusive.

Note: initially, s=1 and e=array length

func paraMergeSort (a, s, e):

if (s == e): } Base case.
return a [s]

m = s + floor ((e-s)/2) } Mid-point of s and e.

l = spawn paraMergeSort (a, s, m) } Sort the left side on a
new thread.

r = paraMergeSort (a, m+1, e) } Sort the right side on this
thread.

Sync } Wait for both sorts to finish.

return paraMerge (l, r) } Recombine the sorted halves
with the fast parallel merge.

The recursion depth and span are $\Theta(\log_2 n)$, each of which uses an $\Theta(\log_2^2 n)$ merge.

The overall runtime is $\Theta(\log_2^3 n)$.

PRAM-EREW Binary Fan-in

Input: n -item array, a

Output: cumulative value in $a[1]$ of operation \oplus

Note: $n = 2^k$

works for any binary associative operation ($+$, $*$, min, max, etc.).
To avoid destroying the input, copy to a duplicate array
first in one parallel step.

for h from 1 to k , do: \leftarrow Loop 1

for i from 1 to $n/2^h$, in parallel do: \leftarrow Loop 2

$$a[i] = a[2i] \oplus a[2i-1]$$

Example: $a = [4, 2, 1, 0, 9, 5, 4, 1]$ operation = addition

• Loop 1, $h=1$ Loop 2, $i = 1$ to 4

$$a = [4, 2, 1, 0, 9, 5, 4, 1]$$

$i=1$ $i=2$ $i=3$ $i=4$

$$a = [6, 1, 14, 5, 9, 5, 4, 1]$$

• Loop 1, $h=2$ Loop 2, $i = 1$ to 2

$$a = [6, 1, 14, 5, 9, 5, 4, 1]$$

$i=1$ $i=2$

$$a = [7, 19, 14, 5, 9, 5, 4, 1]$$

• Loop 1, $h=3$ Loop 2, $i = 1$ to 1

$$a = [7, 19, 14, 5, 9, 5, 4, 1]$$

$i=1$

$$a = [26, 19, 14, 5, 9, 5, 4, 1]$$

The runtime is $\Theta(\log_2 n)$, assuming the operation \oplus is an $O(1)$ action.

For loop 1, in the first iteration loop 2 will need $n/2$ processors.
In the second iteration loop 2 will need $n/4$ processors, then $n/8$,
 $n/16$, etc.

Loop 1 →

Loop 2 follows is,

$$[1, 3, 2] \oplus [1, 2, 3] = [1, 1, 0]$$

$$\text{Iteration} = \text{Index} \quad [1, 4, 2, P, 0, 1, 2, 4] = 0 \quad : \text{aligned}$$

$$P + 4 + 1 = 6, \text{ Loop } 1 = 1, \text{ Loop } 2 =$$

$$[1, 3, 2, P, 0, 1, 2, 4] = 0$$

$$[1, 4, 2, P, 3, 1, 2, 4] = 0$$

$$P + 1 = 2, \text{ Loop } 1 = 1, \text{ Loop } 2 =$$

$$[1, 4, 2, P, 2, 1, 2, 4] = 0$$

$$[1, 4, 2, P, 3, 1, P1, 4] = 0$$

$$P + 1 = 3, \text{ Loop } 1 = 1, \text{ Loop } 2 =$$

$$[1, 4, 2, P, 2, 1, P1, 3] = 0$$

$$[1, 4, 2, P, 2, 1, P1, 2S] = 0$$

PRAM-EREW Array Broadcast

Input: empty n -item array, a value x to broadcast, x

Output: $a[i] = x$ for all $1 \leq i \leq n$.

Note: $n = 2^k$

$a[1] = x$

} Copy to the first cell manually.

for h from 0 to $k-1$, do:

for i from 2^h+1 to 2^{h+1} , in parallel do:

} We will work on the "top half" of increasing powers of two. First just 2, then 3 and 4, then 5 to 8, etc.

$a[i] = a[i - 2^h]$

} Copy the value from the corresponding position in the "bottom half" of the current power of 2.

Example:

$$x = 5$$
$$a = [_, _, _, _, _, _, _, _]$$

Init:

$$a = [5, _, _, _, _, _, _, _] *$$

$$\begin{array}{l} h=0 \\ i=2 \text{ to } 2 \end{array}$$

$$a = [5, 5, _, _, _, _, _, _]$$

$$\begin{array}{l} h=1 \\ i=3 \text{ to } 4 \end{array}$$

$$a = [5, 5, 5, 5, _, _, _, _]$$

$$\begin{array}{l} h=2 \\ i=5 \text{ to } 8 \end{array}$$

$$a = [5, 5, 5, 5, 5, 5, 5, 5]$$

The runtime is $\Theta(\log_2 n)$. In each iteration of the outer loop, the inner loop requires 1, 2, 4, 8... $n/2$ processors.

PRAM EREW Array Membership (Partial)

Input: n-item array, a
n-item array, b , containing the value to check, x

Output: $b[i] = c$ iff $a[i] = x$
 $b[i] = \text{inf}$ iff $a[i] \neq x$

for i from 1 to n , in parallel do: } One processor per index

if ($a[i] == b[i]$): } Store i for matching positions.
 $b[i] = c$

else: } Store i for non-matching positions.
 $b[i] = \text{inf}$

Example: In: $a = [2, 5, 8, 4, 6, 5, 1]$

$b = [5, 5, 5, 5, 5, 5, 5]$

Out: $b = [\infty, 2, \infty, \infty, \infty, 6, \infty]$

Execution takes one parallel step (i.e. $\Theta(1)$) and requires n processors.

PRAM CREW Array Membership (Full)

Input: n -item array, a
item to search for, x

Output: $b[1] = \text{lowest index of } x \text{ in } a, \text{ or infinity.}$

Note: $n = 2^k$

- ① {
 $b = \text{empty } n\text{-item array}$ } This will store n copies of x , for the membership check.
 $\text{Array Broadcast}(b, x)$ } Create n copies of x .
- ② {
 $\text{Array Membership Partial}(a, b)$ } Now b will contain the matching indexes in a .
- ③ {
 $\text{Array Binary Fan In}(b, \min)$ } Find the smallest value in b , which is written to $b[1]$.

Example: Init: $a = [0, 1, 4, 1]$
 $x = 1$

① $b = [1, 1, 1, 1]$

② $b = [\infty, 2, \infty, 4]$

③ $b = [\infty, 2, \infty, 4]$

$b = [2, 4, \infty, 4]$

$b = [2, 4, \infty, 4]$

Answer

PRAM EREW Array Prefix Sums

Input: n values in $a[n \dots 2n-1]$ of array a

Eg. 4 values in $a = [-, -, -, 1, 2, 3, 4]$

Output: prefix sums in $b[n \dots 2n]$

Eg. $b = [-, -, -, 1, 3, 6, 10]$

Note: $n = 2^k$

① { for h from $k-1$ to 0, do:

for i from 2^h to $2^{h+1}-1$, in parallel do:

$$a[i] = a[2i] + a[2i+1]$$

} Move up through each level of the tree, starting above the leaves.

} Make each value the sum of its children. Root will be total sum.

② { $b = \text{empty array, same size as } a$

$$b[1] = a[1]$$

} Manually copy the root.

③ { for h from 1 to k , do:

for i from 2^h to $2^{h+1}-1$, in parallel do:

if (i is even):

$$b[i] = b[i/2] - a[i+1]$$

} Move back down the layers of the tree, starting at the second and ending at the leaves.

} Left-child value = parent - right sibling.

else:

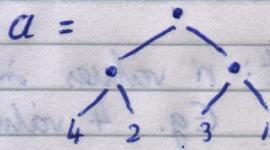
$$b[i] = b[(i-1)/2]$$

} Right-child value = parent

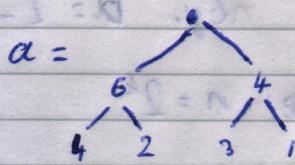
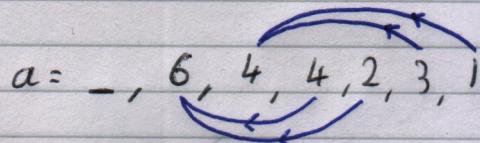
The runtime is $\Theta(\log_2 n)$.

Example

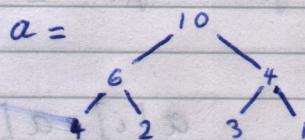
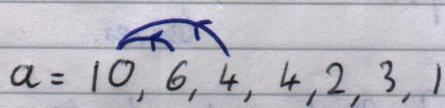
$$a = \underline{\quad}, \underline{\quad}, 4, 2, 3, 1$$



① $h=1, i=2 \text{ to } 3$



② $h=0, i=1 \text{ to } 1$

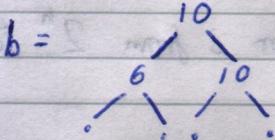


③ $b = 10, \underline{\quad}, \underline{\quad}, \underline{\quad}, \underline{\quad}, \underline{\quad}, \underline{\quad}$

③ $h=1, i=2 \text{ to } 3$

$$a = 10, 6, 4, 4, 2, 3, 1$$

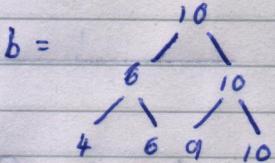
$$b = 10, 6, 10, \underline{\quad}, \underline{\quad}, \underline{\quad}, \underline{\quad}$$



③ $h=2, i=4 \text{ to } 7$

$$a = 10, 6, 4, 4, 2, 3, 1$$

$$b = 10, 6, 10, 4, 6, 9, 10$$



Result: $b = 10, 6, 10, \underline{4, 6, 9, 10}$

PRAM CR*W List Ranking

($\log \max$)

Input: parent array of n -item linked list, p

Output: distance array, d , such that $d[i]$ gives i 's distance from the root

Note: $n = 2^k$

① { for i from 1 to n , do: } Start the root at distance = 0,
 if ($p[i] == i$): all others at distance = 1.
 $d[i] = 0$
 else:
 $d[i] = 1$

for h from 1 to k , do: } $k = \log_2(n)$, which is the maximum
 number of steps needed to point every
 node to the root by parent pointer-
 jumping.

for i from 1 to n , in parallel do: } At every node...

② if ($p[i] := p[p[i]]$): } If it hasn't been pointer-jumped to
 the root yet...

③ $d[i] = d[i] + d[p[i]]$ } Add the distance of the
④ $p[i] = p[p[i]]$ parent, then "jump past" the
 parent.

Runtime is $\Theta(\log n)$, requiring $O(n)$ processors at each iteration.

Example:

$$p = 5 \ 8 \ 1 \ 6 \ 2 \ 3 \ 7 \ 7$$

$$d = 1 \ 1 \ 1 \ 1 \ 1 \ 1 \ 0 \ 1 \ 0$$

$$h=1, i = 1 \ 2 \ 3 \ 4 \ 5 \ 6 \ 7 \ 8 \quad l = n = 8$$

$$pp = 2 \ 7 \ 5 \ 3 \ 8 \ 1 \ 7 \ 7 \ 3 \ ②$$

$$d = 2 \ 2 \ 2 \ 2 \ 2 \ 2 \ 0 \ 1 \ 0 \ ③$$

$$p = 2 \ 7 \ 5 \ 3 \ 8 \ 1 \ 7 \ 7 \ 3 \ ④$$

$$h=2, i = 1 \ 2 \ 3 \ 4 \ 5 \ 6 \ 7 \ 8$$

$$pp = 7 \ 7 \ 8 \ 5 \ 7 \ 2 \ 7 \ 7 \ ②$$

$$d = 4 \ 2 \ 4 \ 4 \ 3 \ 4 \ 0 \ 1 \ ③$$

$$p = 7 \ 7 \ 8 \ 5 \ 7 \ 2 \ 7 \ 7 \ ④$$

$$h=3, i = 1 \ 2 \ 3 \ 4 \ 5 \ 6 \ 7 \ 8$$

$$pp = 7 \ 7 \ 7 \ 7 \ 7 \ 7 \ 7 \ 7 \ ②$$

$$d = 4 \ 2 \ 5 \ 7 \ 3 \ 6 \ 0 \ 1 \ ③$$

↑ Answer.

PRAM CR&W Root Finding

shallow

Input: forest of 1 or more trees stored in the parent array p

Output: roots, r , for all nodes.

Note: $n = 2^k$

for i from 1 to n , in parallel do:
 for h from 1 to k , do:
 $p[i] = p[p[c]]$

} Use pointer jumping to flatten
all trees.

$$r[c] = p[i]$$

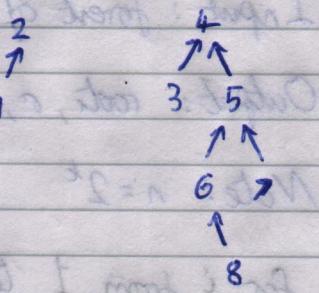
} Store the roots for all nodes.

Runtime is $\Theta(\log_2 n)$, requiring $O(n)$ processors at each iteration.

Example

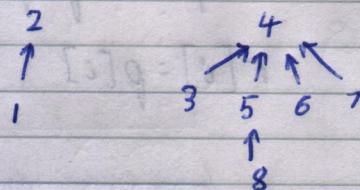
Init: Every cell in bubble sort is moved to front of array

$$p = 2 \ 2 \ 4 \ 4 \ 4 \ 5 \ 5 \ 6$$



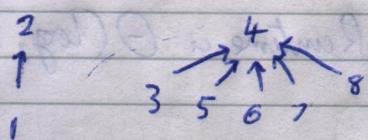
Loop 1:

$$p = 2 \ 2 \ 4 \ 4 \ 4 \ 4 \ 4 \ 5$$



Loop 2:

$$p = 2 \ 2 \ 4 \ 4 \ 4 \ 4 \ 4 \ 4$$



Loop 3:

No change

Output:

$$r = 2 \ 2 \ 4 \ 4 \ 4 \ 4 \ 4 \ 4$$

PRAM CR*W Matrix * Vector Multiplication

symmetric

Input: $n \times n$ -matrix, A
 n -vector, B

$$\begin{bmatrix} 2 \\ 1 \\ 2 \\ 2 \end{bmatrix} = B \quad \begin{bmatrix} 1 & 4 & 2 & 1 \\ 1 & 0 & 4 & 3 \\ 2 & 1 & 2 & 8 \\ 2 & 0 & 1 & 2 \end{bmatrix} = A$$

Output: n -vector, $Q = A \times B$

Note: $n = 2^k$

① { for r, c from 1 to n , in parallel do: } For each vector value, multiply it by each value in the corresponding column of the matrix and store the result in the temporary matrix T .

$$T[r, c] = A[r, c] \times B[c]$$

{ for r from 1 to n , in parallel do: } For each row in parallel:

② { for h from 1 to k , do:
 for c from 1 to $n/2^k$, in parallel do:
 $T[r, c] = T[r, 2c] + T[r, 2c-1]$ } Add up the row using binary fan-in.

③ { Q = empty n -vector
 for i from 1 to n , in parallel do:
 $Q[i] = T[i, 1]$ } Copy the sums to the output vector.

The runtime is $\Theta(\log_2 n)$ and the algorithm requires $\Theta(n^2)$ processors.

Example

$$A = \begin{bmatrix} 1 & 2 & 4 & 1 \\ 5 & 4 & 0 & 1 \\ 8 & 2 & 1 & 2 \\ 2 & 1 & 9 & 3 \end{bmatrix} \quad B = \begin{bmatrix} 2 \\ 1 \\ 3 \\ 2 \end{bmatrix}$$

A 4×4 matrix : Input I
B 4×1 matrix : Input II
 $B \times A = Q$ Output - n rows = Output I

$$\textcircled{1} \quad T = \begin{bmatrix} 1 \times 2 & 2 \times 1 & 4 \times 3 & 1 \times 2 \\ 5 \times 2 & 4 \times 1 & 0 \times 3 & 1 \times 2 \\ 8 \times 2 & 2 \times 1 & 1 \times 3 & 2 \times 2 \\ 2 \times 2 & 1 \times 1 & 9 \times 3 & 3 \times 2 \end{bmatrix} = \begin{bmatrix} 2 & 2 & 12 & 2 \\ 10 & 4 & 0 & 2 \\ 16 & 2 & 3 & 4 \\ 4 & 1 & 27 & 6 \end{bmatrix}$$

Output : 4×4 matrix : Output II } \textcircled{1}

$$\textcircled{2} \quad h=1 \quad T = \begin{bmatrix} 4 & 14 & 12 & 2 \\ 14 & 2 & 0 & 2 \\ 18 & 7 & 3 & 4 \\ 5 & 33 & 27 & 6 \end{bmatrix}$$

Output : 4×4 matrix : Output III } \textcircled{2}

$$\textcircled{2} \quad h=2 \quad T = \begin{bmatrix} 18 & 14 & 12 & 2 \\ 16 & 2 & 0 & 2 \\ 25 & 7 & 3 & 4 \\ 38 & 33 & 27 & 6 \end{bmatrix}$$

Output : 4×4 matrix : Output IV } \textcircled{3}

$$\textcircled{3} \quad Q = \begin{bmatrix} 18 \\ 16 \\ 25 \\ 38 \end{bmatrix}$$

Output : 4×1 matrix = Output V } \textcircled{3}

PRAM CR*W Matrix*Matrix Multiplication

Input: $n \times n$ -matrices, A and B

Output: $n \times n$ -matrix, $Q = A \times B$

Note: $n = 2^k$

① { for r, c, p from 1 to n , in parallel do: } Each cell in the $n \times n$ output matrix is the sum of n multiplication. This step does n^3 multiplications in parallel, storing results in the temporary 3D array T . Think of T as a 2D array of lists, where the list at $T[r, c]$ is the list of values that will be summed for position (r, c) in the output matrix. $T[r, c, p]$ is position p in the "list" at $T[r, c]$.

$$T[r, c, p] = A[r, p] \times B[p, c]$$

for r, c from 1 to n , in parallel do: } For each output matrix position...
 for h from 1 to k , do
 for p from 1 to $\frac{n}{2^h}$, in parallel do: } Add up the "list" for each cell using binary sum in.

$$T[r, c, p] = T[r, c, 2p] + T[r, c, 2p-1]$$

$Q = \text{empty } n \times n \text{ matrix}$

for r, c from 1 to n , in parallel do:

$Q[r, c] = T[r, c, 1]$

} Copy the sums to the output matrix.

The runtime is $\Theta(n \log_2 n)$ with n^3 processors.

Example:

$$A = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} \quad B = \begin{bmatrix} 5 & 6 \\ 7 & 8 \end{bmatrix}$$

① $T = \begin{bmatrix} [1 \times 5, 2 \times 7] & [1 \times 6, 2 \times 8] \\ [3 \times 5, 4 \times 7] & [3 \times 6, 4 \times 8] \end{bmatrix} = \begin{bmatrix} [5, 14] & [6, 16] \\ [15, 28] & [18, 32] \end{bmatrix}$

② $T = \begin{bmatrix} [19, 14] & [22, 16] \\ [43, 28] & [50, 32] \end{bmatrix}$

③ $Q = \begin{bmatrix} 19 & 22 \\ 43 & 50 \end{bmatrix}$

PRAM CRCW-Priority Array Sum

Input: n -item array, a

Output: sum in $a[1]$

$s = 0$

} Initial sum

for i from 1 to n , in parallel do:

$s = s + a[i]$

} Compute the sum in
one parallel step

$a[1] = s$

} Write output.

The runtime is $\Theta(1)$ with $\Theta(n)$ processors.

PRAM CRCW Array Minimum/Maximum

Input: n -item array, a

Output: min/max of a in $a[1:n]$

Note: $n = 2^k$

$m = \text{empty } n\text{-item array}$
for i from 1 to n , in parallel do:
 $m[i] = 0$

} m will store 1 for all
"losers" and a single zero
for the "winner".

* > for max
< for min

for i, j from 1 to n , in parallel do:
if ($a[i] > a[j]$):
 $m[j] = 1$

} Compare all pairs, writing 1s
for the "losers".

for i, j from 1 to n , in parallel do:
if ($m[i] = 0$ and $i < j$):
 $m[j] = 1$

} Compare all pairs again,
writing 1s for all positions
higher than each "winner". This
keeps only the lowest-index
"winner".

for i from 1 to n , in parallel do:
if ($m[i] = 0$):
Result = $a[i]$
IndexOfResult = i

} keep the remaining "winner".

The runtime is $\Theta(1)$, requiring $\Theta(n)$ processors.

The CW model does not matter, because only "1" is ever
written. Priority, arbitrary and common would all work.

PRAM CRCW Array Membership

Input: n -item array, a
item to check, x

Output: (Cw-priority) lowest index containing x , or inf
(Cw-arbitrary) some index containing x , or inf

Index = inf . } Initial state: not found.

for i from 1 to n , in parallel do: } Check every index in parallel.

if $[a[i] == x]$: } Store a matching index
Index = i

Contains = (Index != inf) } This gives a boolean value that
works for Cw-priority and
Cw-arbitrary.

The runtime is $\Theta(1)$, requiring $\Theta(n)$ processes to do $\Theta(n)$ parallel steps.

This algorithm does not work with Cw-common.

PRAM CRCW Array Sorting

Input: n -item array, a

Output: n -item array, b , containing the sorted contents of a

$w = \text{empty } n\text{-item array}$

for i from 1 to n , in parallel do:

$w[i] = 1$

for i, j from 1 to n , in parallel do:

if ($i < j$):

if ($a[i] \leq a[j]$):

$w[j] = w[j] + 1$

else:

$w[i] = w[i] + 1$

for i from 1 to n , in parallel do:

$b[w[i]] = a[i]$

} w will store, for each position i , the number of items smaller than or equal to $a[i]$. Everything starts with one equal: itself.

} For every position, compare all higher positions. If $a[i]$ is less than or equal to $a[j]$, increase j 's score in w . Otherwise, increase i 's score.

The algorithm requires $O(n^2)$ processors and runs in $\Theta(1)$.

Example

$$a = 2 \ 1 \ 4 \ 6 \ 2 \ 8$$

Initial setup: Step I

$$w = 1 \ 1 \ 1 \ 1 \ 1 \ 1$$

$$w = 2 \ 1 \ 2 \ 2 \ 2 \ 2$$

} $i=1, j=2 \text{ to } 6$
 $w[1]=2$ is now fixed

$$w = 2 \ 1 \ 3 \ 3 \ 3 \ 3$$

} $i=2, j=3 \text{ to } 6$
 $w[2]=1$ is now fixed

$$w = 2 \ 1 \ 4 \ 4 \ 3 \ 4$$

} $i=3, j=4 \text{ to } 6$
 $w[3]=4$ is now fixed

$$w = 2 \ 1 \ 4 \ 5 \ 3 \ 5$$

} $i=4, j=5 \text{ to } 6$
 $w[4]=5$ is now fixed

$$w = 2 \ 1 \ 4 \ 5 \ 3 \ 6$$

} $i=5, j=6 \text{ to } 6$
 $w[5]=3$ is now fixed

Output phase:

$$\begin{array}{l} i \\ \hline 1 \rightarrow b[w[1]] = a[1] \rightarrow b[2] = 2 \\ 2 \rightarrow b[w[2]] = a[2] \rightarrow b[1] = 1 \\ 3 \rightarrow b[w[3]] = a[3] \rightarrow b[4] = 4 \\ 4 \rightarrow b[w[4]] = a[4] \rightarrow b[5] = 6 \\ 5 \rightarrow b[w[5]] = a[5] \rightarrow b[3] = 2 \\ 6 \rightarrow b[w[6]] = a[6] \rightarrow b[6] = 8 \end{array} \left. \right\} b = 1 \ 2 \ 2 \ 4 \ 6 \ 8$$

(10) is even but removing (10) gives nothing off

1D Mesh Even/Odd Transposition Sort

Model: 1D mesh, M , with n processors

Input: n values, held as variable x on processors 0 to $n-1$

Output: n values, sorted, held on the same mesh.

for s from 0 to $n-1$; do:

if (s is even):

for each even processor label i , in parallel do:
compareExchange($M(i).x, M(i+1).x$)

} On even cycles, even processors are compared with their right neighbour.

else:

for each odd processor label i , in parallel do:
compareExchange($M(i).x, M(i+1).x$)

} On odd cycles, odd processors are compared with their right neighbour.

func compareExchange(p, q) } Put two items in ascending order.
if ($q < p$):
swap(p, q)

The routine is $O(n)$, because n steps are needed to guarantee that a value can be shuffled along the length of the mesh (e.g. worst case scenario of a reverse sorted array).

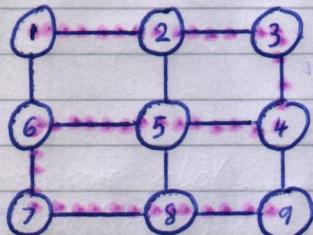
2D Mesh Snake Order Sort

Model: 2D mesh, M , with $q \times q$ processors.

Input: $n = q \times q$ values stored in the mesh M .

Output: mesh values sorted in snake order.

Snake Order:



for s from 0 to $\text{ceil}(\log_2 n)$, do:

if (s is even):

for all rows, in parallel do:
sort row with 1D mesh sort

else:

for all columns, in parallel do:
sort column with 1D mesh sort

The runtime is $\Theta(q \log_2 n)$, because each 1D mesh sort takes $\Theta(q)$.

2D Mesh Broadcast

Model: 2D mesh, M , with $n \times m$ processors. $\xrightarrow{m \uparrow^n}$

Input: value to broadcast at $M(0,0).x$

Output: $M(r, c).x = M(0,0).x$ for all $0 \leq r < n$, $0 \leq c < m$.

for c from 0 to $n-1$, do: } Spread value across the top row,
 $M(0, c+1).x = M(0, c)$ left to right.

for c from 0 to $n-1$, in parallel do: } For every column in parallel...

for r from 0 to $m-1$, do: } Spread the value down each column,
 $M(r+1, c) = M(r, c)$ top to bottom.

The runtime is $\Theta(n+m)$.

2D Mesh Array Min/Max

Model: 2D mesh, M , with $n \times m$ processors. 

Input: nm values stored at variable xe on processors in the mesh.

Output: min/max of the nm values.

for c from 0 to $n-1$, in parallel do: } for each column in parallel...

for r from $m-1$ to 0, do:

$$M(r, c).Temp = M(r+1, c).xe$$

$$M(r, c).xe = \min/\max(Temp, xe)$$

} Bring the min/max to the top by starting one up from the bottom, reading the node below, then storing the min/max, until the top row is reached.

for c from $n-1$ to 0, do:

$$M(0, c).Temp = M(0, c+1).xe$$

$$M(0, c).xe = \min/\max(Temp, xe)$$

} Work along the top row in the same fashion, from right to left, to bring the min/max into the top-left processor.

return $\overset{m}{\overbrace{M(0, 0)}}.xe$

The runtime is $\Theta(n+m)$.

2D Mesh Membership

Model: 2D mesh, M , with $n \times m$ processors. $m \xrightarrow{n}$

Input: $n \cdot m$ values stored in the mesh at memory variable x .
value to search for, y .

Output: Lowest index of processor containing y , or inf .

2d-mesh-broadcast y into $M.y$ for all processors } Give each processor a
copy of y to read.

for r from 0 to $n-1$, in parallel do: } for each processor in
parallel...

for c from 0 to $m-1$, in parallel do:

if $(M(r,c).x == M(r,c).y)$: } Store the index of matching processors.

$$M(r,c).s = r \times n + c$$

else:

$$M(r,c).s = \text{inf}$$

} Store inf for non-matches.

2d-mesh-array-min on $M.s$ for all processors } find the smallest s
value.

return $M(0,0).s$

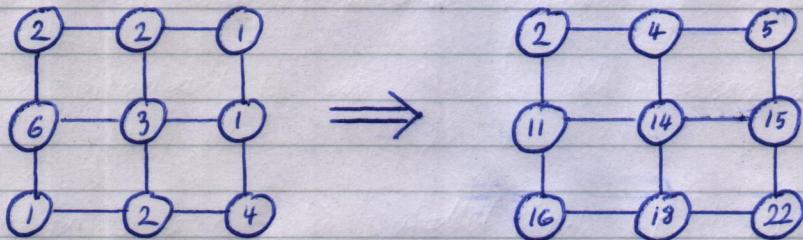
The runtime is $\Theta(n+m)$.

2D Mesh Array Prefix Sums

Model: 2D mesh, M , with $n \times m$ processors. $m \uparrow^n$

Input: nm values stored as variable x in the mesh.

Output: prefix sums of the variables:



for r from 0 to $m-1$, in parallel do: } For each row in parallel:

$$M(r, 0).s = M(r, 0).x$$

for c from 1 to $n-1$, do

$$M(r, c).s = M(r, c).x + M(r, c-1).s$$

} Add up each row, left to right, so that $1 \rightarrow 2 \rightarrow 3$ becomes $1 \rightarrow 3 \rightarrow 6$.

for r from 1 to $m-1$, do:

$$M(r, n-1).s = M(r, n-1) + M(r-1, n-1)$$

} Add up the last column, top to bottom, in the same fashion. "n-1" just gives the index of the last column.

for r from 1 to $m-1$, in parallel do:

for c from 0 to $n-2$, in parallel do:

$$M(r, c).s = M(r, c).s + M(r-1, n-1)$$

} On every row except the top row, for every cell except the rightmost, add the value of the rightmost cell from the row above.

The runtime is $\Theta(n+m)$.

Directed Ring Leader Election

Model: directed ring of message passing nodes.

Output: one node has status = "leader"
all other nodes have status = "not leader"

Note: the algorithm is as viewed at any node, n.

status = "unknown"

send ID to neighbour

} Initially no one knows their status, so everyone sends their ID as it might be the max.

while status = "unknown"

on receipt of message:

} Until a node learns its status, it listens for messages.

if (message == "leader"):

 status = "not leader"

 send "leader" to neighbour

} A leader was found elsewhere, so update status and pass the message on.

if (message is an ID):

 if (message.id == ID):

 status = "leader"

 send "leader" to neighbour

} If a node receives its own ID then it is the leader, so it updates status and starts telling everyone.

if (message.id > ID)

 send message.id to neighbour

} Drop lower IDs, but keep sending higher IDs around the ring.

$O(n)$ messages will be sent.

General Network Leader Election

Model: network of message-passing nodes.

Output: one node has status = "leader"
all other nodes have status = "not leader"

Note: algorithm is as viewed at any node, n
all nodes know the network diameter, d.

$n.\text{maxId} = n.\text{id}$ } To start, each node assumes it has the max ID.

for r from 1 to d , do: } Repeat "diameter" times.

send $n.\text{maxId}$ to all neighbours } Send the current max ID...

for all IDs received, i : } Then overwrite it if a higher
if ($i > n.\text{maxId}$): ID is received.
 $n.\text{maxId} = i$

if ($n.\text{maxId} == n.\text{id}$): } A node is the leader iff it maintains
status = "leader" its max ID after d
else: cycles of "send and overwrite."
status = "not leader"

d is the longest time required for the max ID to spread through the entire network.

General Network BFS Tree Finding

Model: network of message-forwarding nodes.

Output: every node knows its parent and children, forming a tree.

Note: algorithm is started from a specific node (such as the leader)

At the start node(s):

for all nodes, n :
 $n.\text{status} = \text{"not in tree"}$
 $n.\text{parent} = \text{"unknown"}$
 $n.\text{children} = \{\}$

} Initialise a default state for all nodes.

$s.\text{parent} = s$
 $s.\text{status} = \text{"in tree"}$

} The start node, s , declares itself to be the root.

send "join tree" to all neighbours } This starts the algorithm below.

At every node (n):

on receipt of "join tree" message:
if ($n.\text{status} == \text{"not in tree"}$):
select $v = \text{lowest ID of "join tree" sender}$
 $n.\text{parent} = v$

} When told to join a tree, pick the lowest ID of parents saying "join tree" if not already in one, then join that tree.

send "your child" to v
send "not your child" to all neighbours except v

} Send child status messages to all neighbours.

send "join tree" to all neighbours } Continue spreading the tree.

on receipt of "your child" message from v :
add v to $n.\text{children}$

} Record children as they arrive.

when child status is known from all neighbours except parent:

if ($n.\text{children} == \{\}$):
send "terminated" to parent

else:
when "terminated" is received from all children:
send "terminated" to parent.

} Tell the parent we are done when all neighbours have child status and all children (if any) have finished.

$O(e)$ messages will be sent, where e is the number of edges.

arbeit gutvom - gegen 30 Minuten : Jahr M

sen 5 giorni, molti hanno già usato le loro penne: fuga!

so that I also suggest a work before it will go : still
before 30

(2) Show that $\exists A$

The right things to prioritize:
when {
 "and neither" = either or
 "immediate" = handle it
 { } = prioritize it

5 Next number is always first off if
from left end "last in" = newest, 2

reduced ultimate life time of materials due to 'end wear' has

(A) down strong FA

where $\omega = \text{SI form} = \text{rads}$

using this last { v of 'high way' here
method to & never { v shows nothing the & 'high way' has
seen at previous visits) { evidence it is "working" here
well no further basis) { v with easier 'high way' to figure no
evidence at v like

so that it is not : two types consider the most usual is what kind when
the new and one
kind and another
the two relate
(one if) variable
but only one
variable the most common is "backward" new
known as "backward" new

Sequential Graph Components

Input: list of nodes, each knowing their outgoing edges.

Output: list of sets of nodes, each being a component.

Note: NOT PARALLEL! Included only for comparison.

n = list of nodes

s = list of sets

while (n is not empty):

$c = \text{lowest-index node in } n$
 $s[c] = \{c\}$
 $n. \text{remove}(c)$

} Pick the lowest-index node, create a set for it, then remove it from the "working set".

do:

fan out from C with BFS
add found nodes to $s[c]$
remove found nodes from n
while (BFS finds more nodes)

} Discover c 's entire component using repeated BFS until no more nodes are found.

return s

The runtime is $\Theta(|\text{Nodes}| + |\text{Edges}|)$, which can be close to $O(|\text{Nodes}|^2)$ for densely connected graphs.

Parallel Graph Components

Model: graph structure (vertices V and edges E)

Output: each vertex will know its super-vertex (the lowest label within its component).

Note: $S(v) = v$'s super-vertex
 $N(v) = \text{neighbours of } v$

for each v in V , in parallel do: } Each vertex starts as its own super-vertex.

$S(v) = v$ } $\log_2(|V|)$ cycles needed for complete solution.

for each v in V , in parallel do:

$$\min NS = \min \{S(j) \text{ for each } j \in N(v)\}$$

$$S(v) = \min(v, \min NS)$$

} Find the minimum of the S-values held by all neighbours of each vertex, then overwrite each node's S-value if it is bigger than the found minimum.

for each v in V , in parallel do:

repeat $\log_2(|V|)$ times:

$$S(v) = S(S(v))$$

} Find the S-value root for all vertices.

The runtime is $\Theta(\log_2^2 |V|)$.

Parallel Minimum Spanning Tree (Boruvka's Method)

Model: graph structure

Output: MST for graph

Note: unique edge weights are required

Each vertex starts as its own single-member super-vertex.

While edges remain:

- Select the min weight edge out of each super-vertex
- Use selected edges to merge super-vertices
- Remove non-selected edges within super-vertices (self-loops)
- If multiple edges exist between components, remove all but the cheapest.
- Add all selected edges to the MST.